

Column-Oriented Database Systems


VLDB
2009
Tutorial



Part 1: Stavros Harizopoulos (HP Labs)

Part 2: Daniel Abadi (Yale)

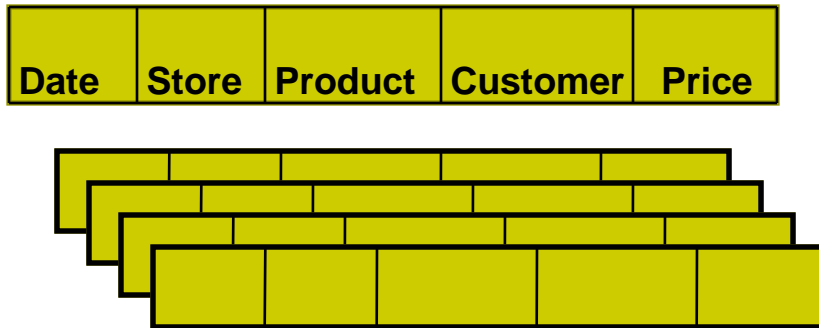
Part 3: Peter Boncz (CWI)






What is a column-store?

row-store



- + easy to add/modify a record
- might read in unnecessary data

column-store



- + only need to read in relevant data
- tuple writes require multiple accesses

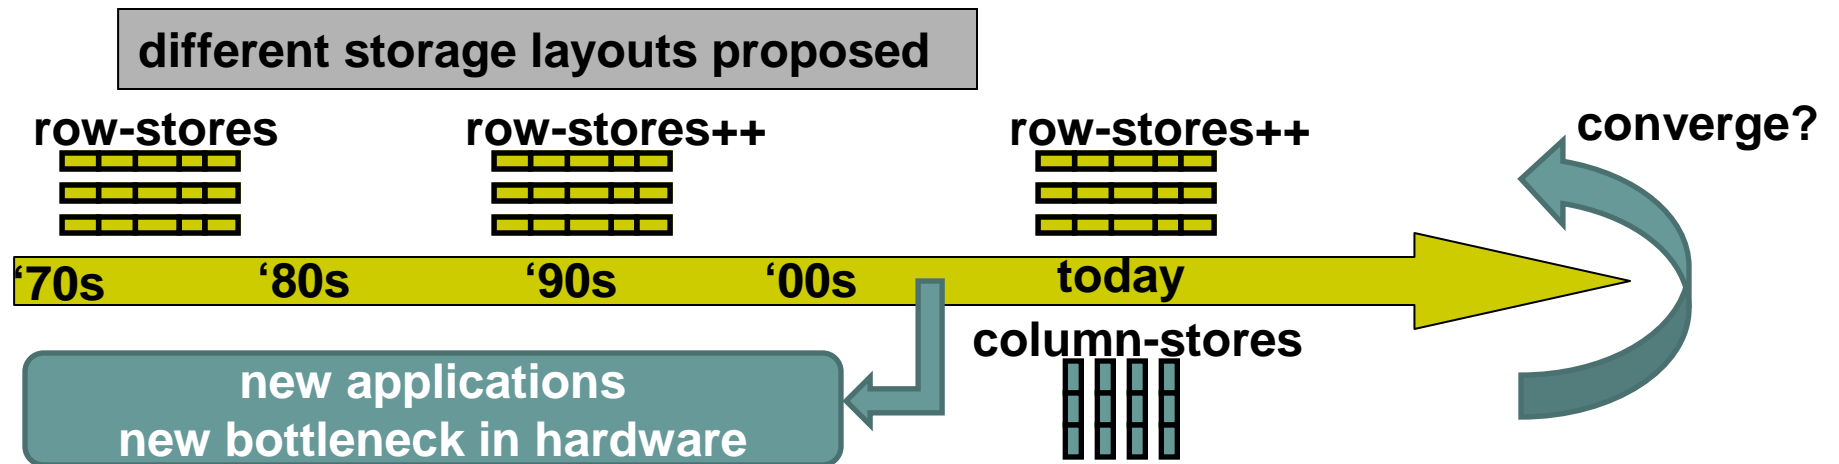
=> suitable for read-mostly, read-intensive, large data repositories





Are these two fundamentally different?

- 1 The only fundamental difference is the storage layout
- 1 However: we need to look at the big picture



- 1 How did we get here, and where we are heading
- 1 What are the column-specific optimizations?
- 1 How do we improve CPU efficiency when operating on Cs

Part 1

Part 2

Part 3





Outline

- 1 Part 1: Basic concepts — *Stavros*
 - 1 Introduction to key features
 - 1 From DSM to column-stores and performance tradeoffs
 - 1 Column-store architecture overview
 - 1 Will rows and columns ever converge?
- 1 Part 2: Column-oriented execution — *Daniel*
- 1 Part 3: MonetDB/X100 and CPU efficiency — *Peter*






Telco Data Warehousing example

1 Typical DW installation

1 Real-world example

“One Size Fits All? - Part 2: Benchmarking Results” Stonebraker et al. CIDR 2007



QUERY 2

```
SELECT account.account_number,  
sum (usage.toll_airtime),  
sum (usage.toll_price)  
FROM usage, toll, source, account  
WHERE usage.toll_id = toll.toll_id  
AND usage.source_id = source.source_id  
AND usage.account_id = account.account_id  
AND toll.type_ind in ('AE', 'AA')  
AND usage.toll_price > 0  
AND source.type != 'CIBER'  
AND toll.rating_method = 'IS'  
AND usage.invoice_date = 20051013  
GROUP BY account.account_number
```

	Column-store	Row-store
Query 1	2.06	300
Query 2	2.20	300
Query 3	0.09	300
Query 4	5.24	300
Query 5	2.88	300


Why? Three main factors (next slides)



Telco example explained (1/3): *read efficiency*



row store




read pages containing entire rows

one row = 212 columns!

is this typical? (it depends)

What about vertical partitioning?
(it does not work with ad-hoc queries)

column store



read only columns needed

in this example: 7 columns

caveats:

- “select * ” not any faster
- clever disk prefetching
- clever tuple reconstruction





Telco example explained (2/3): *compression efficiency*

- 1 Columns compress better than rows
 - 1 Typical row-store compression ratio 1 : 3
 - 1 Column-store 1 : 10
- 1 Why?
 - 1 Rows contain values from different domains
=> more entropy, difficult to dense-pack
 - 1 Columns exhibit significantly less entropy
 - 1 Examples:

Male, Female, Female, Female, Male
1998, 1998, 1999, 1999, 1999, 2000
 - 1 Caveat: CPU cost (use lightweight compression)



Telco example explained (3/3): *sorting & indexing efficiency*



- 1 Compression and dense-packing free up space
 - 1 Use multiple overlapping column collections
 - 1 Sorted columns compress better
 - 1 Range queries are faster
 - 1 Use sparse clustered indexes

**What about heavily-indexed row-stores?
(works well for single column access,
cross-column joins become increasingly expensive)**





Additional opportunities for column-stores

- 1 Block-tuple / vectorized processing
 - 1 Easier to build block-tuple operators
 - 1 Amortizes function-call cost, improves CPU cache performance
 - 1 Easier to apply vectorized primitives
 - 1 Software-based: bitwise operations
 - 1 Hardware-based: SIMD
- 1 Opportunities with compressed columns
 - 1 *Avoid* decompression: operate directly on compressed
 - 1 *Delay* decompression (and tuple reconstruction)
 - 1 Also known as: *late materialization*
- 1 Exploit columnar storage in other DBMS components
 - 1 Physical design (both static and dynamic)

Part 3


more
in Part 2

See: *Database Cracking*, from CWI



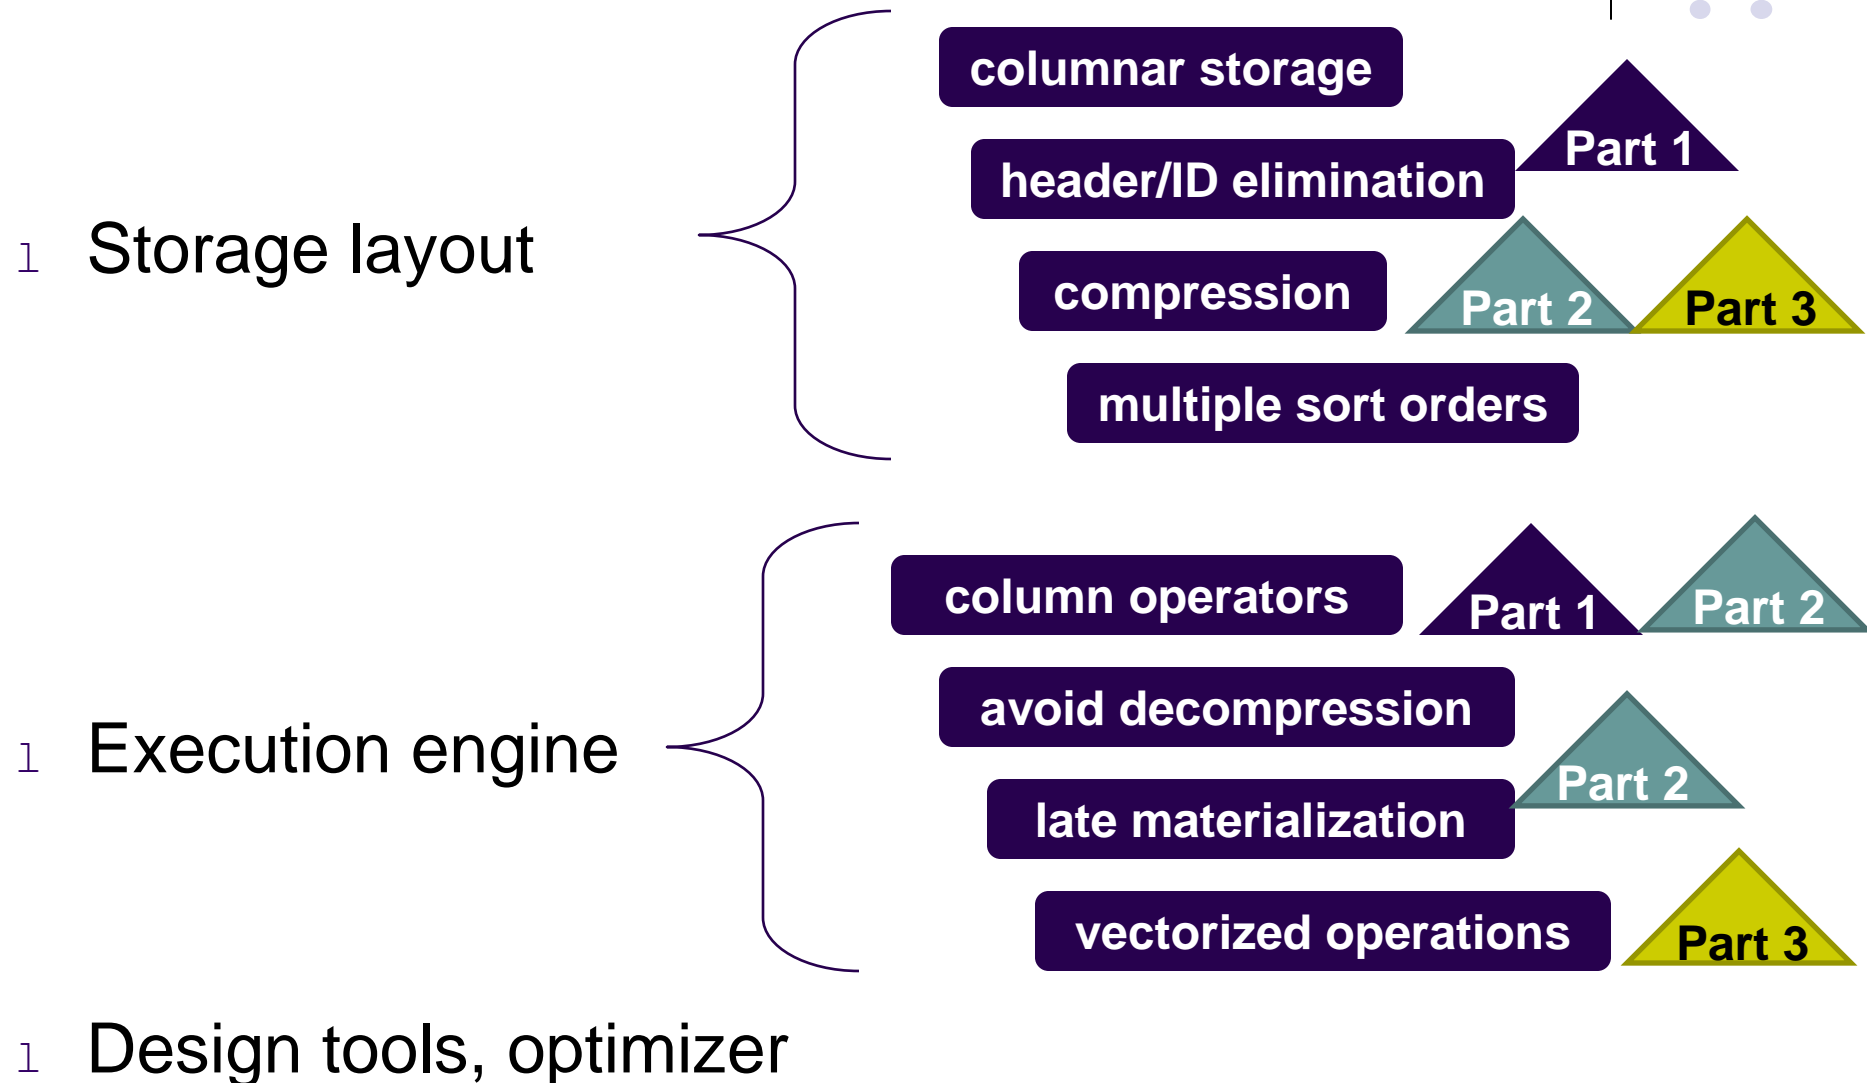
“Column-Stores vs Row-Stores: How Different are They Really?” Abadi, Hachem, and Madden. SIGMOD 2008.

Effect on C-Store performance





Summary of column-store key features





Outline

1 Part 1: Basic concepts — *Stavros*

1 Introduction to key features

1 From DSM to column-stores and performance tradeoffs

1 Column-store architecture overview

1 Will rows and columns ever converge?

1 Part 2: Column-oriented execution — *Daniel*

1 Part 3: MonetDB/X100 and CPU efficiency — *Peter*





From DSM to Column-stores

70s -1985:

TOD: Time Oriented Database – Wiederhold et al.
"A Modular, Self-Describing Clinical Databank System," *Computers and Biomedical Research*, 1975
More 1970s: Transposed files, Lorie, Batory,

"An overview of cantor: a new system for data analysis"
Karasalo, Svensson, SSDBM 1983

1985: DSM paper

"A decomposition storage model"
Copeland and Khoshafian. SIGMOD 1985.

1990s: Commercialization through SybaseIQ

Late 90s – 2000s: Focus on main-memory performance

1 DSM "on steroids" [1997 – now] CWI: MonetDB

1 Hybrid DSM/NSM [2001 – 2004] Wisconsin: PAX, Fractured Mirrors

Michigan: Data Morphing

CMU: Clotho

2005 – : Re-birth of read-optimized DSM as "column-store"

MIT: C-Store

CWI: MonetDB/X100

10+ startups

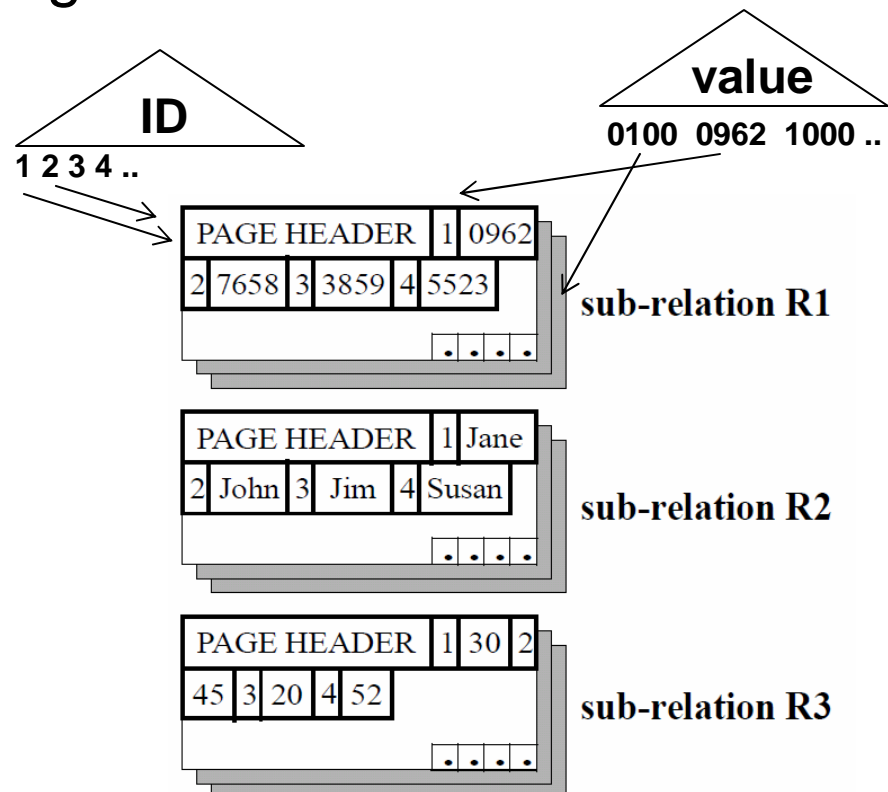





The original DSM paper

“A decomposition storage model” Copeland and Khoshafian. SIGMOD 1985.

- 1 Proposed as an alternative to NSM
- 1 2 indexes: clustered on ID, non-clustered on value
- 1 Speeds up queries projecting few columns
- 1 Requires more storage





Memory wall and PAX

1 90s: Cache-conscious research

from: “Cache Conscious Algorithms for Relational Query Processing.”
Shatdal, Kant, Naughton. VLDB 1994.

to: “Database Architecture Optimized for the New Bottleneck: Memory Access.”
Boncz, Manegold, Kersten. VLDB 1999.

and: “DBMSs on a modern processor: Where does time go?” Ailamaki, DeWitt, Hill, Wood. VLDB 1999.

1 PAX: Partition Attributes Across

1 Retains NSM I/O pattern

1 Optimizes cache-to-RAM communication

“Weaving Relations for Cache Performance.”
Ailamaki, DeWitt, Hill, Skounakis, VLDB 2001.

PAX PAGE

PAGE HEADER		0962	7658
3859	5523		
Jane	John	Jim	Susan
30	52	45	20





More hybrid NSM/DSM schemes

1 Dynamic PAX: Data Morphing


“Data morphing: an adaptive, cache-conscious storage technique.” Hankins, Patel, VLDB 2003.

1 Clotho: custom layout using scatter-gather I/O

“Clotho: Decoupling Memory Page Layout from Storage Organization.” Shao, Schindler, Schlosser, Ailamaki, and Ganger. VLDB 2004.

1 Fractured mirrors

1 Smart mirroring with both NSM/DSM copies



“A Case For Fractured Mirrors.” Ramamurthy, DeWitt, Su, VLDB 2002.





MonetDB (more in Part 3)

- 1 Late 1990s, CWI: Boncz, Manegold, and Kersten
- 1 Motivation:
 - 1 Main-memory
 - 1 Improve computational efficiency by avoiding expression interpreter
 - 1 DSM with virtual IDs natural choice
 - 1 Developed new query execution algebra
- 1 Initial contributions:
 - 1 Pointed out memory-wall in DBMSs
 - 1 Cache-conscious projections and joins
 - 1 ...





2005: the (re)birth of column-stores

- 1 New hardware and application realities
 - 1 Faster CPUs, larger memories, disk bandwidth limit
 - 1 Multi-terabyte Data Warehouses
- 1 New approach: combine several techniques
 - 1 Read-optimized, fast multi-column access, disk/CPU efficiency, light-weight compression
- 1 C-store paper:
 - 1 First comprehensive design description of a column-store
- 1 MonetDB/X100
 - 1 “proper” disk-based column store
- 1 Explosion of new products





Performance tradeoffs: columns vs. rows

DSM traditionally was not favored by technology trends
How has this changed?

1 Optimized DSM in “Fractured Mirrors,” 2002

1 “Apples-to-apples” comparison

“Performance Tradeoffs in Read-Optimized Databases”
Harizopoulos, Liang, Abadi,
Madden, VLDB’06

1 Follow-up study

“Read-Optimized Databases, In-Depth” Holloway, DeWitt, VLDB’08

1 Main-memory DSM vs. NSM

“DSM vs. NSM: CPU performance tradeoffs in block-oriented query processing” Boncz, Zukowski, Nes, DaMoN’08

1 Flash-disks: a come-back for PAX?

“Fast Scans and Joins Using Flash Drives” Shah, Harizopoulos, Wiener, Graefe. DaMoN’08

“Query Processing Techniques for Solid State Drives”
Tsirogiannis, Harizopoulos, Shah, Wiener, Graefe,
SIGMOD’08






Fractured mirrors: a closer look

- 1 Store DSM relations inside a B-tree
 - 1 Leaf nodes contain values
 - 1 Eliminate IDs, amortize header overhead
 - 1 Custom implementation on Shore

“A Case For Fractured Mirrors” Ramamurthy, DeWitt, Su, VLDB 2002.

Tuple Header	TID	Column Data
	1	a1
	2	a2
	3	a3
	4	a4
	5	a5





Similar: storage density comparable to column stores

“Efficient columnar storage in B-trees” Graefe. Sigmod Record 03/2007.




Fractured mirrors: performance

From PAX paper:



- 1 Chunk-based tuple merging
 - 1 Read in segments of M pages
 - 1 Merge segments in memory
 - 1 Becomes CPU-bound after 5 pages

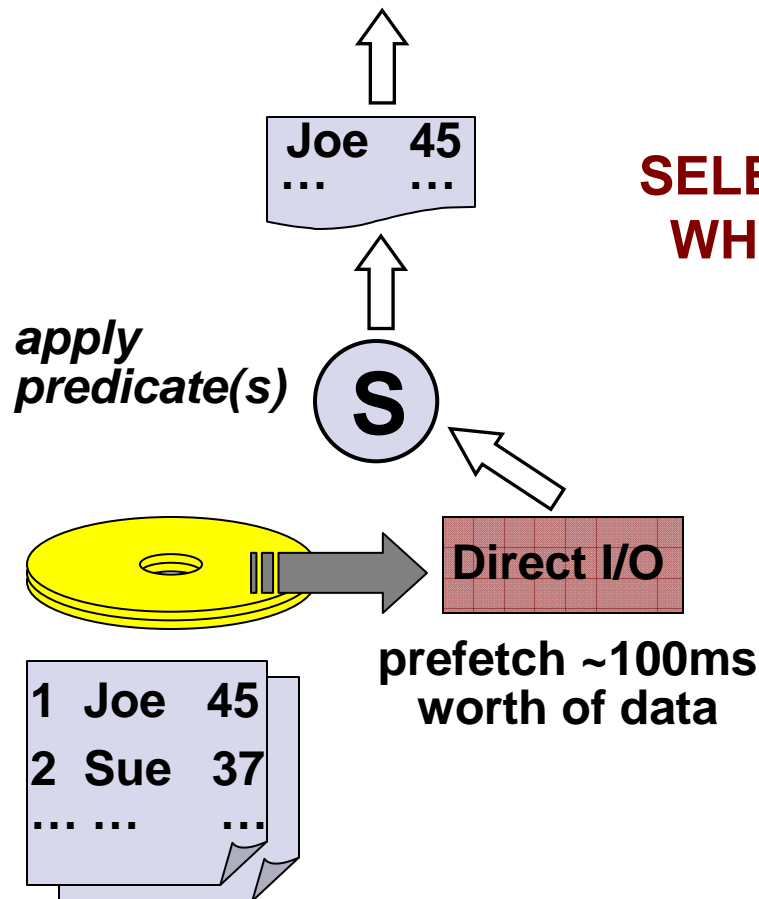


Column-scanner implementation

“Performance Tradeoffs in Read-Optimized Databases”
Harizopoulos, Liang, Abadi,
Madden, VLDB’06




row scanner



SELECT name, age
WHERE age > 40

column scanner






Scan performance

- 1 Large prefetch hides disk seeks in columns
- 1 Column-CPU efficiency with lower selectivity
- 1 Row-CPU suffers from memory stalls
- 1 Memory stalls disappear in narrow tuples
- 1 Compression: similar to narrow

not shown,
details in the paper

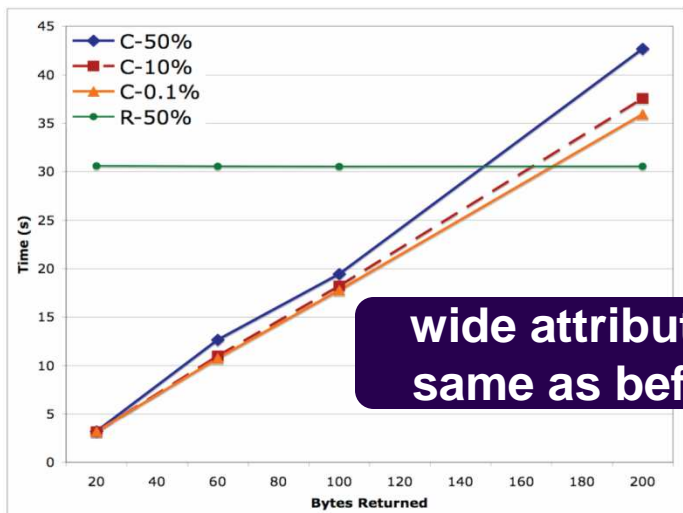




Even more results


- Same engine as before
- Additional findings

“Read-Optimized Databases, In-Depth” Holloway, DeWitt, VLDB’08



**wide attributes:
same as before**

**narrow & compressed tuple:
CPU-bound!**



Non-selective queries, narrow tuples, favor well-compressed rows


Materialized views are a win

Scan times determine early materialized joins

**Column-joins are
covered in part 2!**




“Performance Tradeoffs in Read-Optimized Databases”
Harizopoulos, Liang, Abadi,
Madden, VLDB’06





Varying prefetch size




- 1 No prefetching hurts columns in single scans





Varying prefetch size



- 1 No prefetching hurts columns in single scans
- 1 Under competing traffic, columns outperform rows for any prefetch size



CPU Performance

“DSM vs. NSM: CPU performance trade offs in block-oriented query processing”
Boncz, Zukowski, Nes, DaMoN'08



- 1 Benefit in on-the-fly conversion between NSM and DSM
- 1 DSM: sequential access (block fits in L2), random in L1
- 1 NSM: random access, SIMD for grouped Aggregation

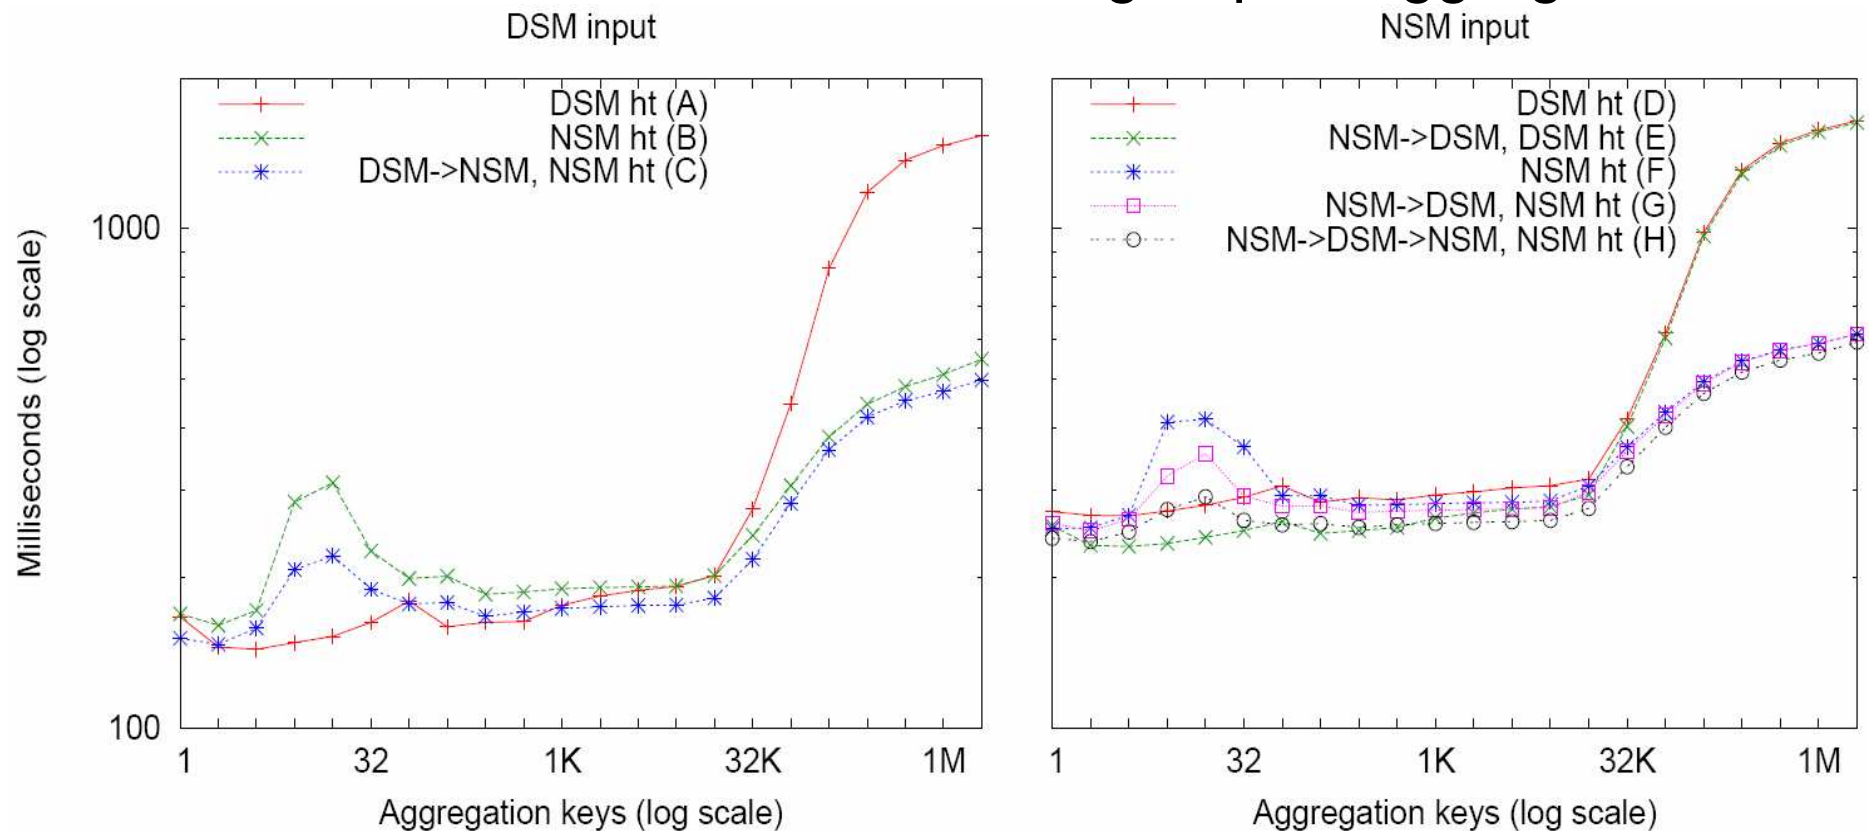


Figure 5: TPC-H Q1, with a varying number of keys and different data organizations (ht – hash table)



New storage technology: Flash SSDs

- 1 Performance characteristics
 - 1 very fast random reads, slow random writes
 - 1 fast sequential reads and writes
 - 1 Price per bit (capacity follows)
 - 1 cheaper than RAM, order of magnitude more expensive than Disk
 - 1 Flash Translation Layer introduces unpredictability
 - 1 avoid random writes!
 - 1 Form factors not ideal yet
 - 1 SSD (⌘ small reads still suffer from SATA overhead/OS limitations)
 - 1 PCI card (⌘ high price, limited expandability)
-
- 1 Boost Sequential I/O in a simple package
 - 1 Flash RAID: very tight bandwidth/cm³ packing (4GB/sec inside the box)
 - 1 Column Store Updates
 - 1 useful for delta structures and logs
 - 1 Random I/O on flash fixes unclustered index access
 - 1 still suboptimal if I/O block size > record size
 - 1 therefore column stores profit much less than horizontal stores
 - 1 Random I/O useful to exploit secondary, tertiary table orderings
 - 1 the larger the data, the deeper clustering one can exploit





Even faster column scans on flash SSDs

1 New-generation SSDs


30K Read IOps, 3K Write Iops
250MB/s Read BW, 200MB/s Write

1 Very fast random reads, slower random writes

1 Fast sequential RW, comparable to HDD arrays

1 No expensive seeks across columns

1 FlashScan and Flashjoin: PAX on SSDs, inside Postgres



“Query Processing Techniques for Solid State Drives” Tsirogiannis, Harizopoulos, Shah, Wiener, Graefe, SIGMOD’09


mini-pages with no qualified attributes are not accessed






Column-scan performance over time

regular DSM (2001)




from 7x slower

column-store (2006)



..to 1.2x slower


Lineitem (TPCH) 1GB



..to 2x slower

and 3x faster!

..to same



SSD Postgres/PAX (2009)





Outline

- 1 Part 1: Basic concepts — *Stavros*
 - 1 Introduction to key features
 - 1 From DSM to column-stores and performance tradeoffs
 - 1 Column-store architecture overview
 - 1 Will rows and columns ever converge?
- 1 Part 2: Column-oriented execution — *Daniel*
- 1 Part 3: MonetDB/X100 and CPU efficiency — *Peter*





Architecture of a column-store

storage layout

- 1 read-optimized: dense-packed, compressed
- 1 organize in extends, batch updates
- 1 multiple sort orders
- 1 sparse indexes

engine

- 1 block-tuple operators
- 1 new access methods
- 1 optimized relational operators

system-level

- 1 system-wide column support
- 1 loading / updates
- 1 scaling through multiple nodes
- 1 transactions / redundancy



“C-Store: A Column-Oriented DBMS.” Stonebraker et al. VLDB 2005.



C-Store

- 1 Compress columns
- 1 No alignment
- 1 Big disk blocks
- 1 Only materialized views (perhaps many)
- 1 Focus on Sorting not indexing
- 1 Data ordered on anything, not just time
- 1 Automatic physical DBMS design
- 1 Optimize for grid computing
- 1 Innovative redundancy
- 1 Xacts – but no need for Mohan
- 1 Column optimizer and executor





C-Store: only materialized views (MVs)

- 1 **Projection** (MV) is some number of columns from a fact table
- 1 Plus columns in a dimension table – with a 1-n join between Fact and Dimension table
- 1 Stored in order of a storage key(s)
- 1 Several may be stored!
- 1 With a **permutation**, if necessary, to map between them
- 1 Table (as the user specified it and sees it) is not stored!
- 1 No secondary indexes (they are a one column sorted MV plus a permutation, if you really want one)

User view:

EMP (name, age, salary, dept)
Dept (dname, floor)

Possible set of MVs:

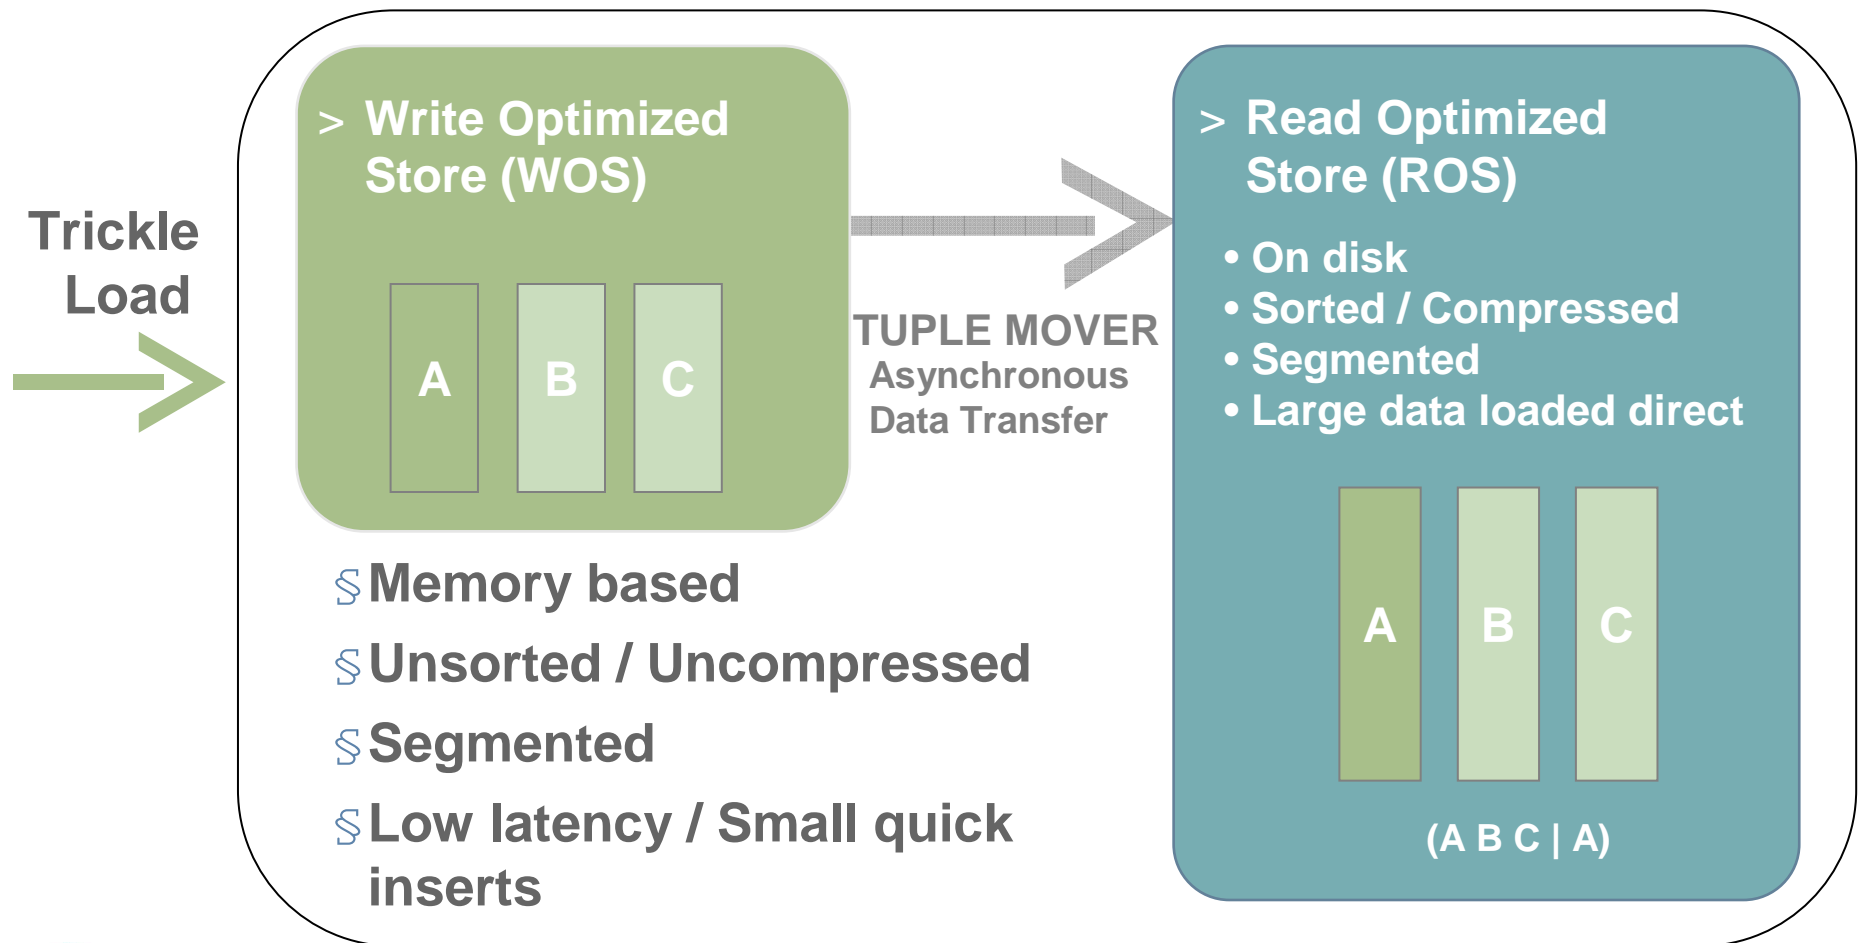
MV-1 (name, dept, floor) in floor order
MV-2 (salary, age) in age order
MV-3 (dname, salary, name) in salary order





Continuous Load and Query (Vertica)

Hybrid Storage Architecture





Loading Data (Vertica)

> INSERT, UPDATE, DELETE


> Bulk and Trickle Loads

§ COPY


§ COPY DIRECT

> User loads data into logical Tables

> Vertica loads atomically into storage



Automatic
Tuple Mover





Applications for column-stores

- 1 Data Warehousing
 - 1 High end (clustering)
 - 1 Mid end/Mass Market
 - 1 Personal Analytics
- 1 Data Mining
 - 1 E.g. Proximity
- 1 Google BigTable
- 1 RDF
 - 1 Semantic web data management
- 1 Information retrieval
 - 1 Terabyte TREC
- 1 Scientific datasets
 - 1 SciDB initiative
 - 1 SLOAN Digital Sky Survey on MonetDB





List of column-store systems

- 1 Cantor (history)
- 1 Sybase IQ
- 1 SenSage (former Addamark Technologies)
- 1 Kdb
- 1 1010data
- 1 MonetDB
- 1 C-Store/Vertica
- 1 X100/VectorWise
- 1 KickFire
- 1 SAP Business Accelerator
- 1 Infobright
- 1 ParAccel
- 1 Exasol





Outline

- 1 Part 1: Basic concepts — *Stavros*
 - 1 Introduction to key features
 - 1 From DSM to column-stores and performance tradeoffs
 - 1 Column-store architecture overview
 - 1 Will rows and columns ever converge?
- 1 Part 2: Column-oriented execution — *Daniel*
- 1 Part 3: MonetDB/X100 and CPU efficiency — *Peter*






Simulate a Column-Store inside a Row-Store

Date	Store	Product	Customer	Price
01/01	BOS	Table	Mesa	\$20
01/01	NYC	Chair	Lutz	\$13
01/01	BOS	Bed	Mudd	\$79

**Option A:
Vertical Partitioning**

Date		Store		Product		Customer		Price	
TID	Value	TID	Value	TID	Value	TID	Value	TID	Value
1	01/01	1	BOS	1	Table	1	Mesa	1	\$20
2	01/01	2	NYC	2	Chair	2	Lutz	2	\$13
3	01/01	3	BOS	3	Bed	3	Mudd	3	\$79

**Option B:
Index Every Column**





Simulate a Column-Store inside a Row-Store

Date	Store	Product	Customer	Price
01/01	BOS	Table	Mesa	\$20
01/01	NYC	Chair	Lutz	\$13
01/01	BOS	Bed	Mudd	\$79



Option A: Vertical Partitioning

Date		
Value	StartPos	Length
01/01	1	3

Can explicitly run-length encode date

“Teaching an Old Elephant New Tricks.”
Bruno, CIDR 2009.

Option B: Index Every Column



...






Experiments

1 Star Schema Benchmark (SSBM)

Adjoined Dimension Column Index (ADC Index) to Improve Star Schema Query Performance”. O’Neil et. al. ICDE 2008.

1 Implemented by professional DBA

1 Original row-store plus 2 column-store simulations on same row-store product



“Column-Stores vs Row-Stores: How Different are They Really?” Abadi, Hachem, and Madden. SIGMOD 2008.





What's Going On? Vertical Partitions

- 1 Vertical partitions in row-stores:
 - 1 Work well when workload is known
 - 1 ..and queries access disjoint sets of columns
 - 1 See automated physical design
- 1 Do not work well as full-columns
 - 1 TupleID overhead significant
 - 1 Excessive joins

Tuple Header	TID	Column Data
	1	
	2	
	3	

Queries touch 3-4 foreign keys in fact table,
1-2 numeric columns

Complete fact table takes up ~4 GB
(compressed)

Vertically partitioned tables take up 0.7-1.1
GB (compressed)

“Column-Stores vs. Row-Stores:
How Different Are They Really?”
Abadi, Madden, and Hachem.
SIGMOD 2008.





What's Going On? All Indexes Case

1 Tuple construction

1 Common type of query:

```
SELECT store_name, SUM(revenue)
FROM Facts, Stores
WHERE fact.store_id = stores.store_id
      AND stores.country = "Canada"
GROUP BY store_name
```

1 Result of lower part of query plan is a set of TIDs that passed all predicates

1 Need to extract SELECT attributes at these TIDs

1 BUT: index maps value to TID

1 You really want to map TID to value (i.e., a vertical partition)

Tuple construction is SLOW





So....

- 1 All indexes approach is a poor way to simulate a column-store
- 1 Problems with vertical partitioning are NOT fundamental
 - 1 Store tuple header in a separate partition
 - 1 Allow virtual TIDs
 - 1 Combine clustered indexes, vertical partitioning
- 1 So can row-stores simulate column-stores?
 - 1 Might be possible, BUT:
 - 1 Need better support for vertical partitioning at the storage layer
 - 1 Need support for column-specific optimizations at the executor level
 - 1 Full integration: buffer pool, transaction manager, ..
 - 1 When will this happen?
 - 1 Most promising features = soon
 - 1 ..unless new technology / new objectives change the game (SSDs, Massively Parallel Platforms, Energy-efficiency)

See Part 2, Part 3
for most promising
features





End of Part 1

- 1 Basic concepts — *Stavros*
 - 1 Introduction to key features
 - 1 From DSM to column-stores and performance tradeoffs
 - 1 Column-store architecture overview
 - 1 Will rows and columns ever converge?
- 1 Part 2: Column-oriented execution — *Daniel*
- 1 Part 3: MonetDB/X100 and CPU efficiency — *Peter*





Part 2 Outline

- 1 Compression
- 1 Tuple Materialization
- 1 Joins



Column-Oriented Database Systems

VLDB
2009
Tutorial




Compression

“Super-Scalar RAM-CPU Cache Compression”
Zukowski, Heman, Nes, Boncz, ICDE’06

“Integrating Compression and Execution in Column-Oriented Database Systems” Abadi, Madden, and Ferreira, SIGMOD ’06

•Query optimization in compressed database systems” Chen, Gehrke, Korn, SIGMOD’01





Compression

- 1 **Trades I/O for CPU**
- 1 **Increased column-store opportunities:**
 - 1 **Higher data value locality in column stores**
 - 1 **Techniques such as run length encoding far more useful**
 - 1 **Can use extra space to store multiple copies of data in different sort orders**





Run-length Encoding

Quarter Product ID Price

Q1	1	5
Q1	1	7
Q1	1	2
Q1	1	9
Q1	1	6
Q1	2	8
Q1	2	5
...
Q2	1	3
Q2	1	8
Q2	1	1
Q2	2	4
...



Quarter

(value, start_pos, run_length)
(Q1, 1, 300)
(Q2, 301, 350)
(Q3, 651, 500)
(Q4, 1151, 600)

Product ID Price

(value, start_pos, run_length)	
(1, 1, 5)	5
(2, 6, 2)	7
...	2
...	9
(1, 301, 3)	6
(2, 304, 1)	8
...	5
...	...
...	3
...	8
...	1
...	4
...	...





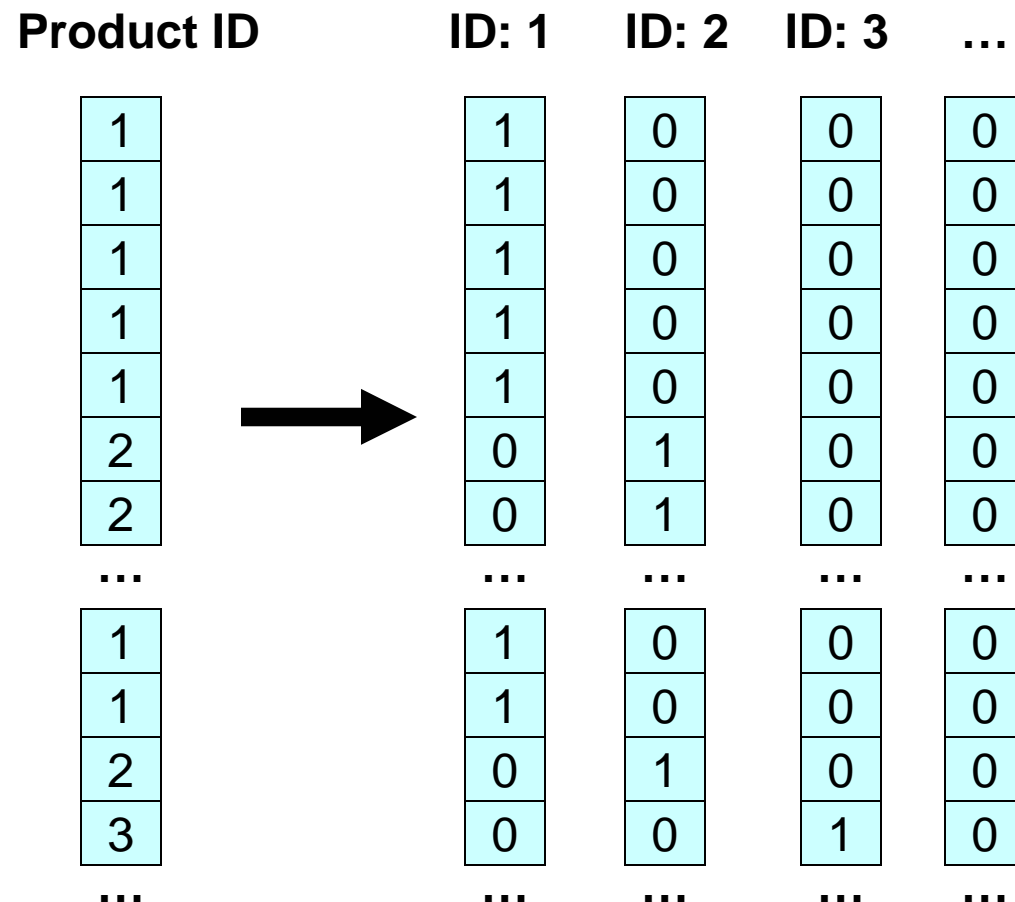
Bit-vector Encoding

- 1 For each unique value, v , in column c , create bit-vector b

- 1 $b[i] = 1$ if $c[i] = v$

- 1 Good for columns with few unique values

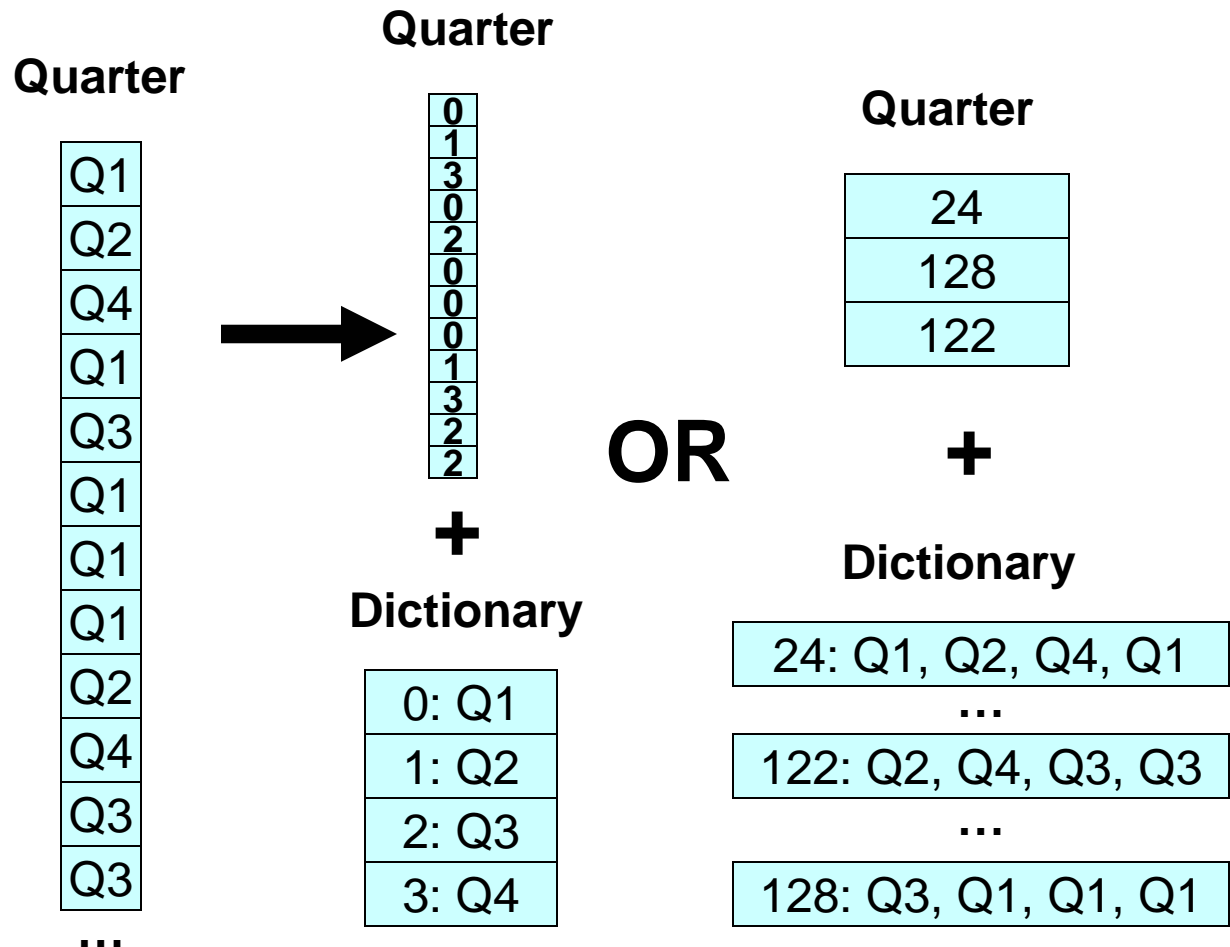
- 1 Each bit-vector can be further compressed if sparse





Dictionary Encoding

- 1 For each unique value create dictionary entry
- 1 Dictionary can be per-block or per-column
- 1 Column-stores have the advantage that dictionary entries may encode multiple values at once

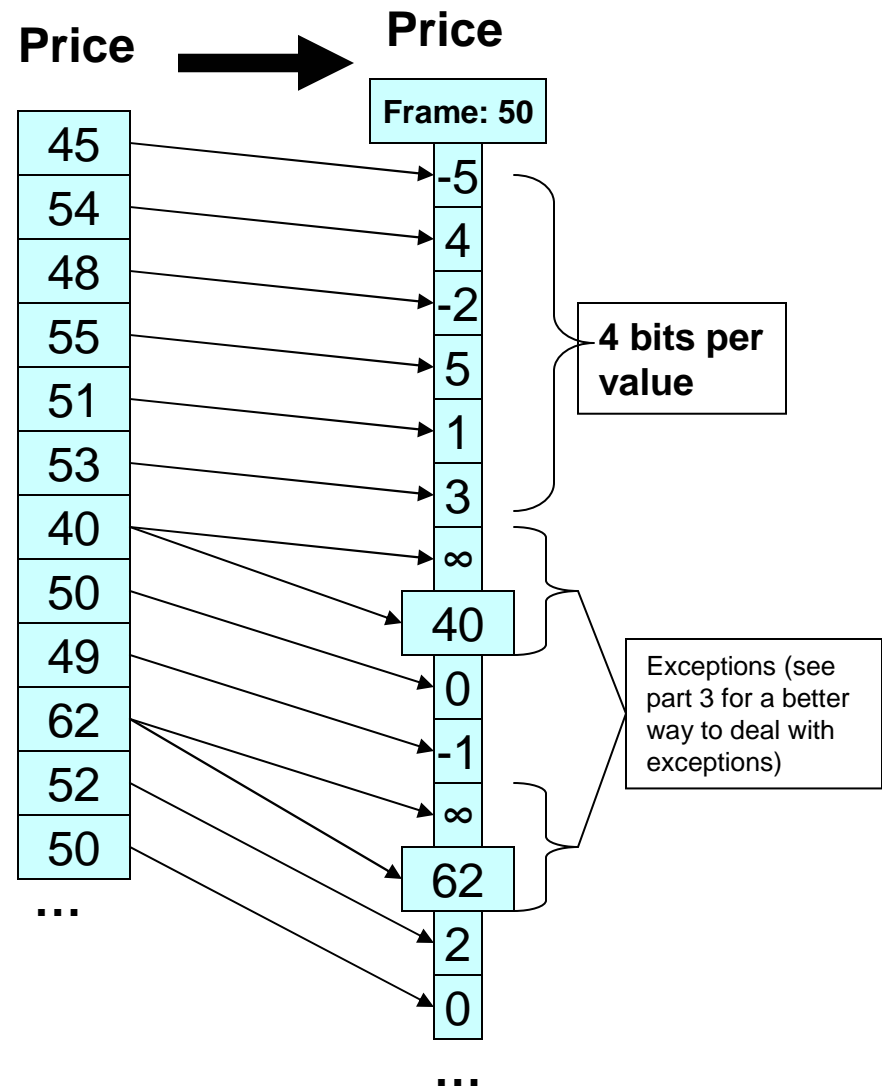




Frame Of Reference Encoding

- 1 Encodes values as b bit offset from chosen frame of reference
- 1 Special escape code (e.g. all bits set to 1) indicates a difference larger than can be stored in b bits
 - 1 After escape code, original (uncompressed) value is written

“Compressing Relations and Indexes” Goldstein, Ramakrishnan, Shaft, ICDE’98






Differential Encoding

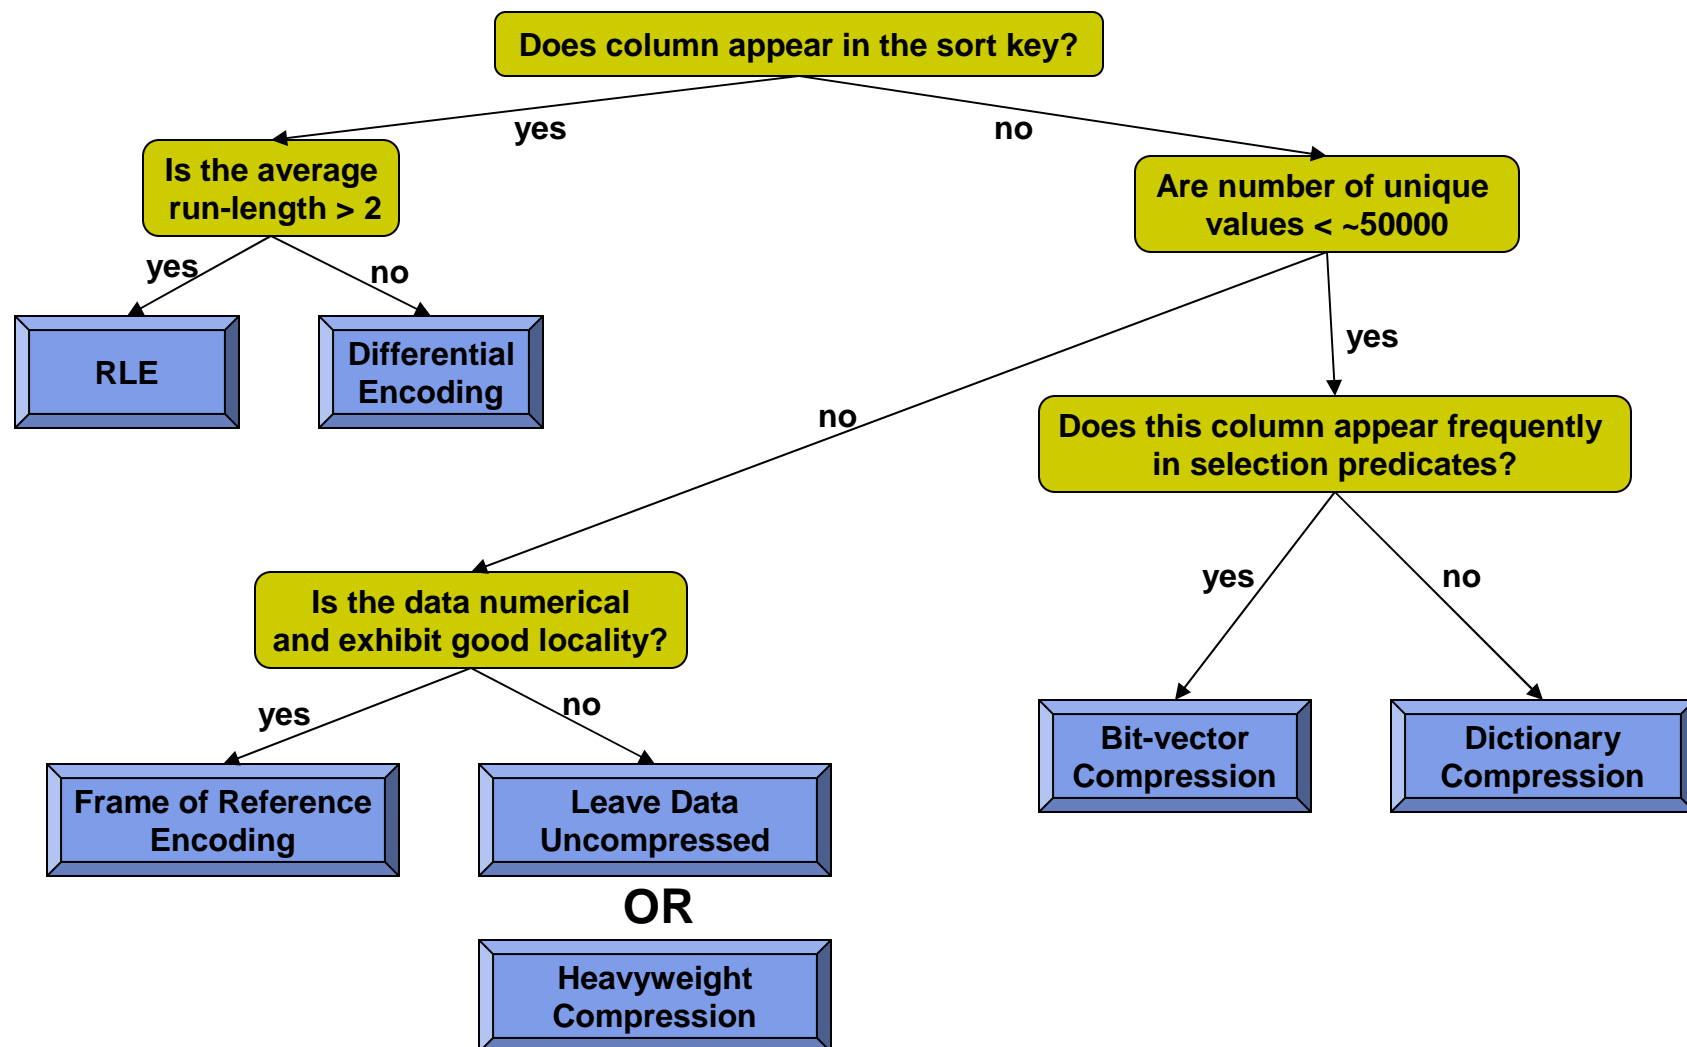
- 1 Encodes values as b bit offset from previous value
- 1 Special escape code (just like frame of reference encoding) indicates a difference larger than can be stored in b bits
 - 1 After escape code, original (uncompressed) value is written
- 1 Performs well on columns containing increasing/decreasing sequences
 - 1 inverted lists
 - 1 timestamps
 - 1 object IDs
 - 1 sorted / clustered columns

“Improved Word-Aligned Binary Compression for Text Indexing”
Ahn, Moffat, TKDE’06





What Compression Scheme To Use?



“Super-Scalar RAM-CPU Cache Compression”
Zukowski, Heman, Nes, Boncz, ICDE’06



Heavy-Weight Compression Schemes

Algorithm	Decompression Bandwidth
BZIP	10 MB/s
ZLIB	80 MB/s
LZO	300 MB/s

- 1 Modern disk arrays can achieve $> 1\text{GB/s}$
- 1 1/3 CPU for decompression $\approx 3\text{GB/s}$ needed
- \approx Lightweight compression schemes are better
- \approx Even better: operate directly on compressed data



“Integrating Compression and Execution in Column-Oriented Database Systems” Abadi et. al, SIGMOD '06



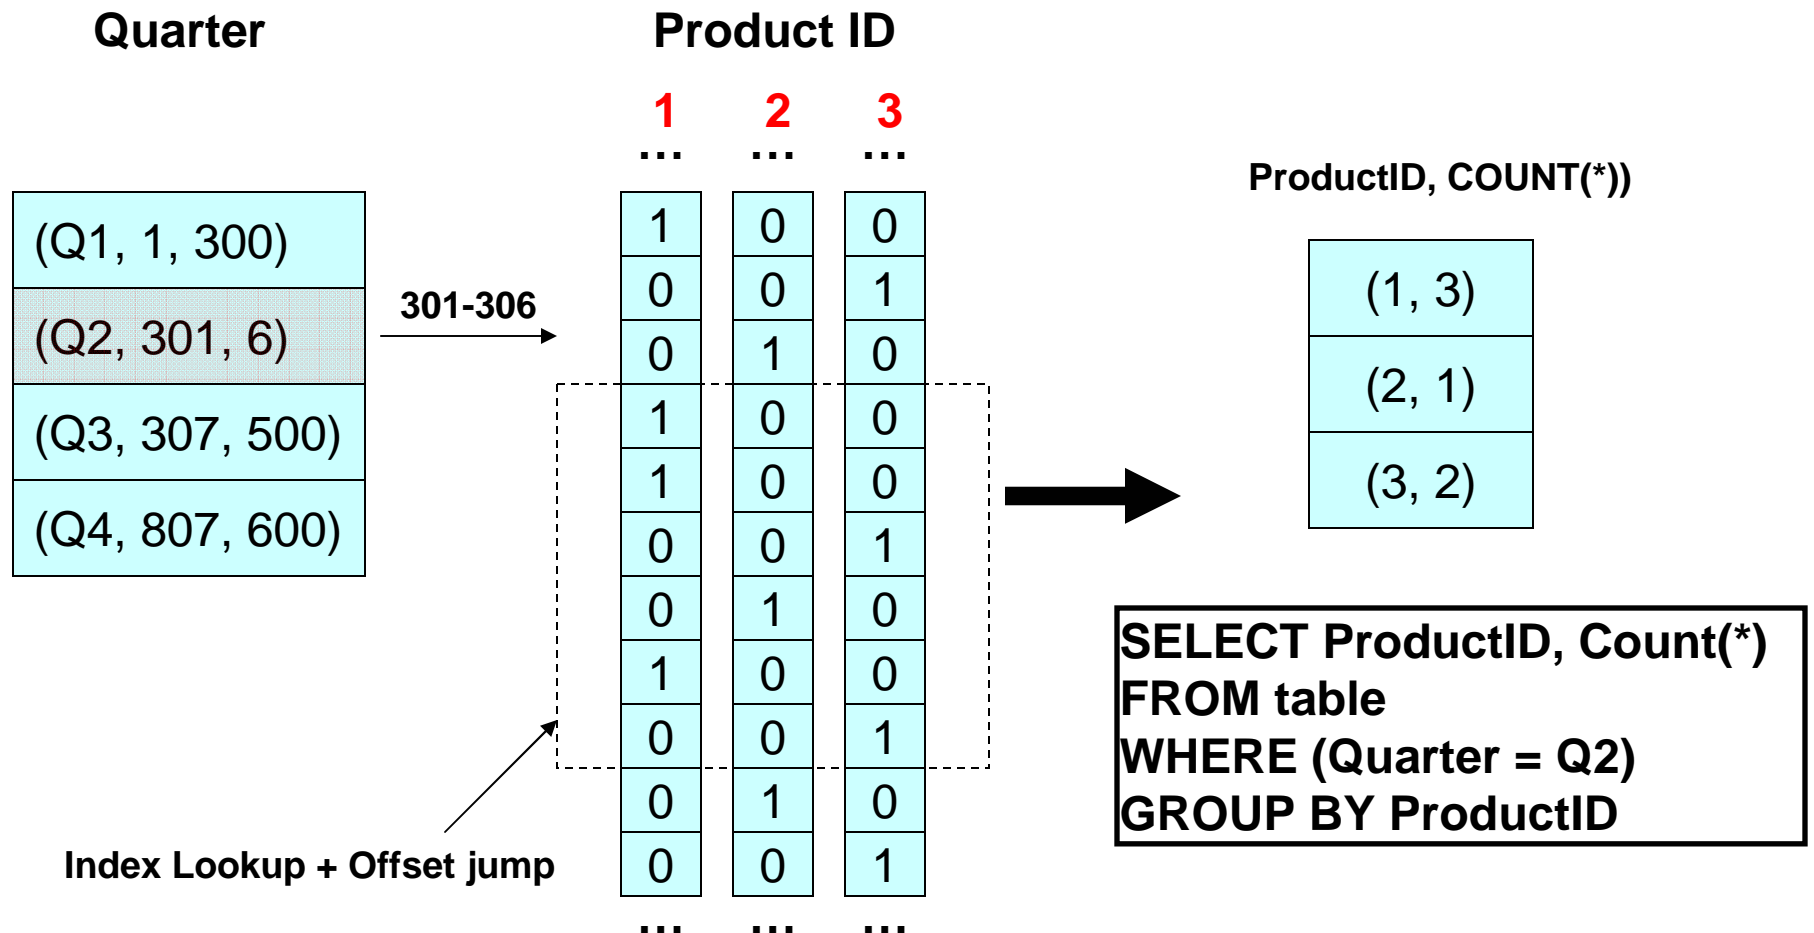
Operating Directly on Compressed Data

- 1 **I/O - CPU tradeoff is no longer a tradeoff**
- 1 **Reduces memory–CPU bandwidth requirements**
- 1 **Opens up possibility of operating on multiple records at once**





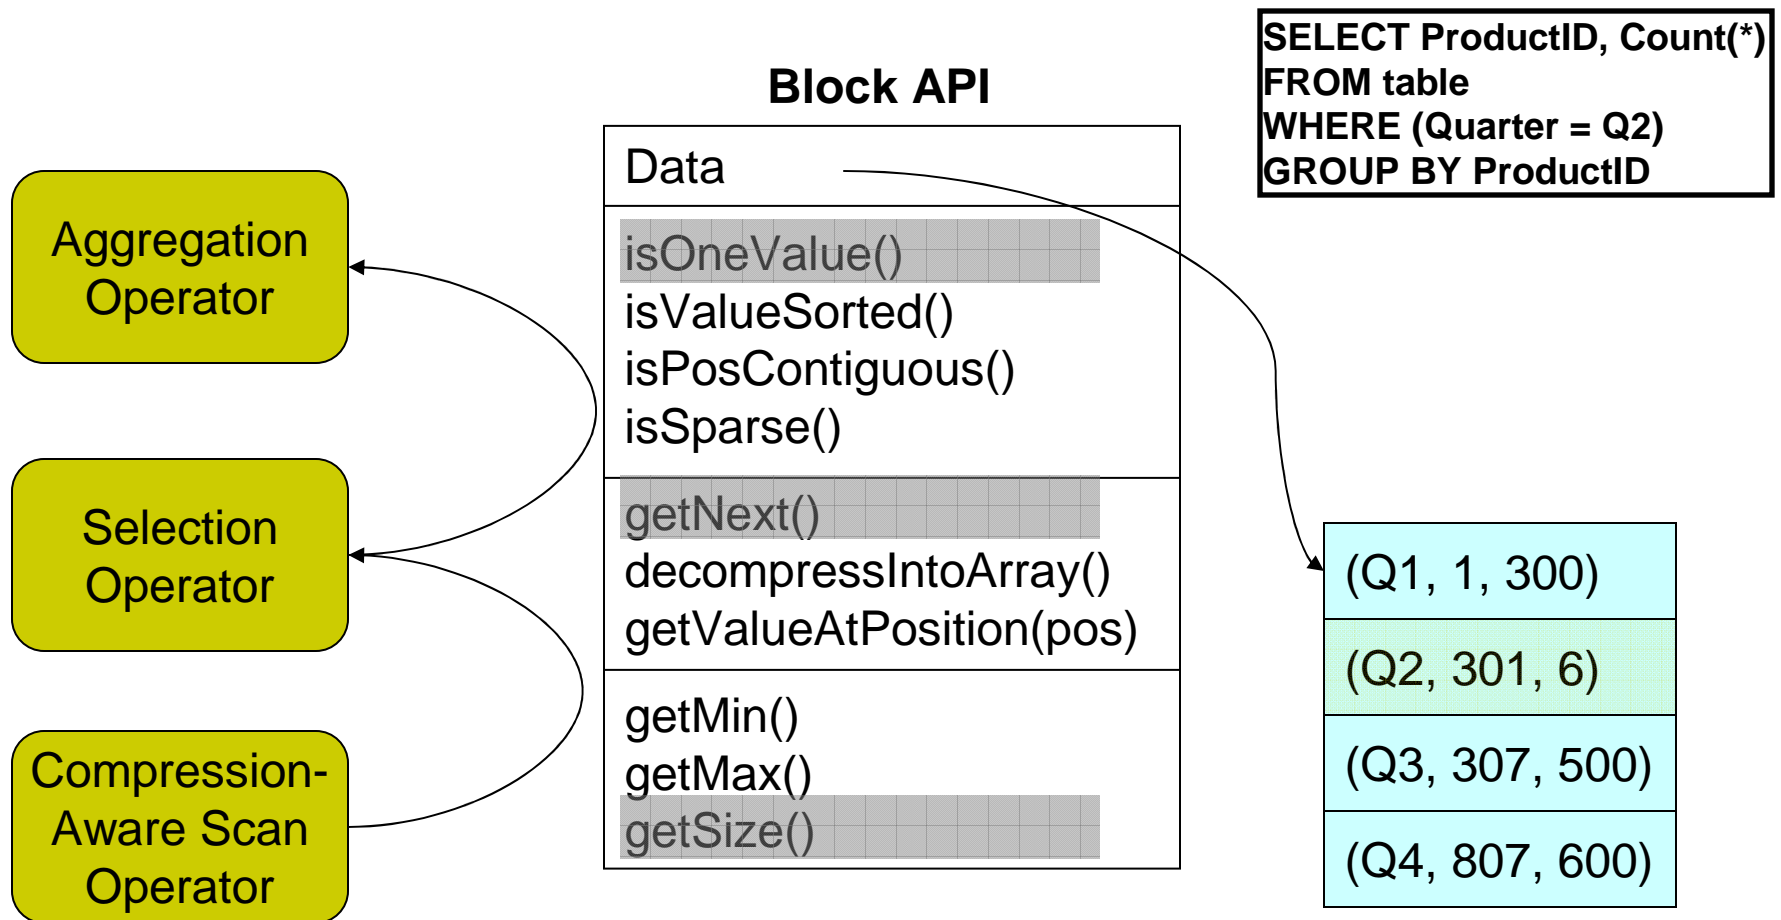
Operating Directly on Compressed Data



“Integrating Compression and Execution in Column-Oriented Database Systems” Abadi et. al, SIGMOD '06



Operating Directly on Compressed Data



Column-Oriented Database Systems

VLDB
2009
Tutorial



Tuple Materialization and Column-Oriented Join Algorithms


“Materialization Strategies in a Column-Oriented DBMS” Abadi, Myers, DeWitt, and Madden. ICDE 2007.

“Self-organizing tuple reconstruction in column-stores”, Idreos, Manegold, Kersten, SIGMOD’09

“Column-Stores vs Row-Stores: How Different are They Really?” Abadi, Madden, and Hachem. SIGMOD 2008.

“Query Processing Techniques for Solid State Drives” Tsirogiannis, Harizopoulos Shah, Wiener, and Graefe. SIGMOD 2009.

“Cache-Conscious Radix-Decluster Projections”, Manegold, Boncz, Nes, VLDB’04






When should columns be projected?

- 1 **Where should column projection operators be placed in a query plan?**
 - 1 **Row-store:**
 - 1 Column projection involves removing unneeded columns from tuples
 - 1 Generally done as early as possible
 - 1 **Column-store:**
 - 1 Operation is almost completely opposite from a row-store
 - 1 Column projection involves reading needed columns from storage and extracting values for a listed set of tuples
 - § This process is called “materialization”
 - 1 **Early materialization: project columns at beginning of query plan**
 - § Straightforward since there is a one-to-one mapping across columns
 - 1 **Late materialization: wait as long as possible for projecting columns**
 - § More complicated since selection and join operators on one column obfuscates mapping to other columns from same table
 - 1 **Most column-stores construct tuples and column projection time**
 - § Many database interfaces expect output in regular tuples (rows)
 - § Rest of discussion will focus on this case





When should tuples be constructed?



QUERY:

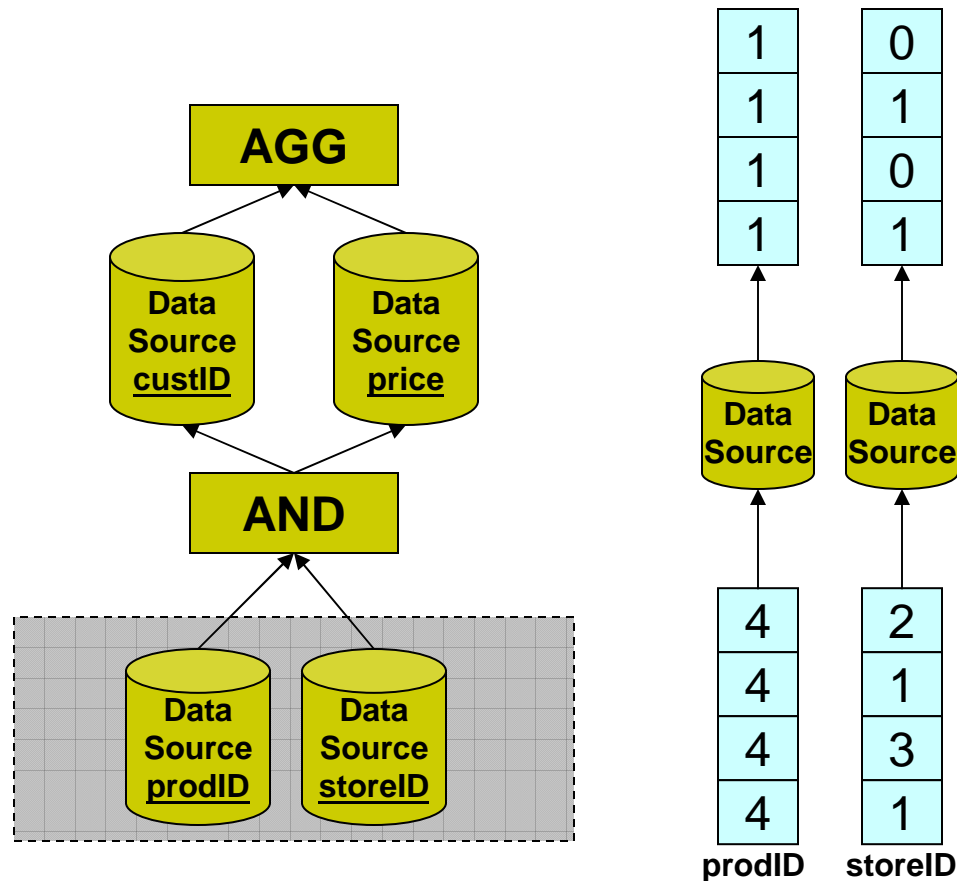
```
SELECT custID, SUM(price)
FROM table
WHERE (prodID = 4) AND
      (storeID = 1) AND
GROUP BY custID
```

- 1 **Solution 1: Create rows first (EM).**
- But:**
 - 1 **Need to construct ALL tuples**
 - 1 **Need to decompress data**
 - 1 **Poor memory bandwidth utilization**





Solution 2: Operate on columns



QUERY:



```
SELECT custID, SUM(price)
FROM table
WHERE (prodID = 4) AND
      (storeID = 1) AND
GROUP BY custID
```

4	2	2	7
4	1	3	13
4	3	3	42
4	1	3	80
prodID	storeID	custID	price





Solution 2: Operate on columns



QUERY:

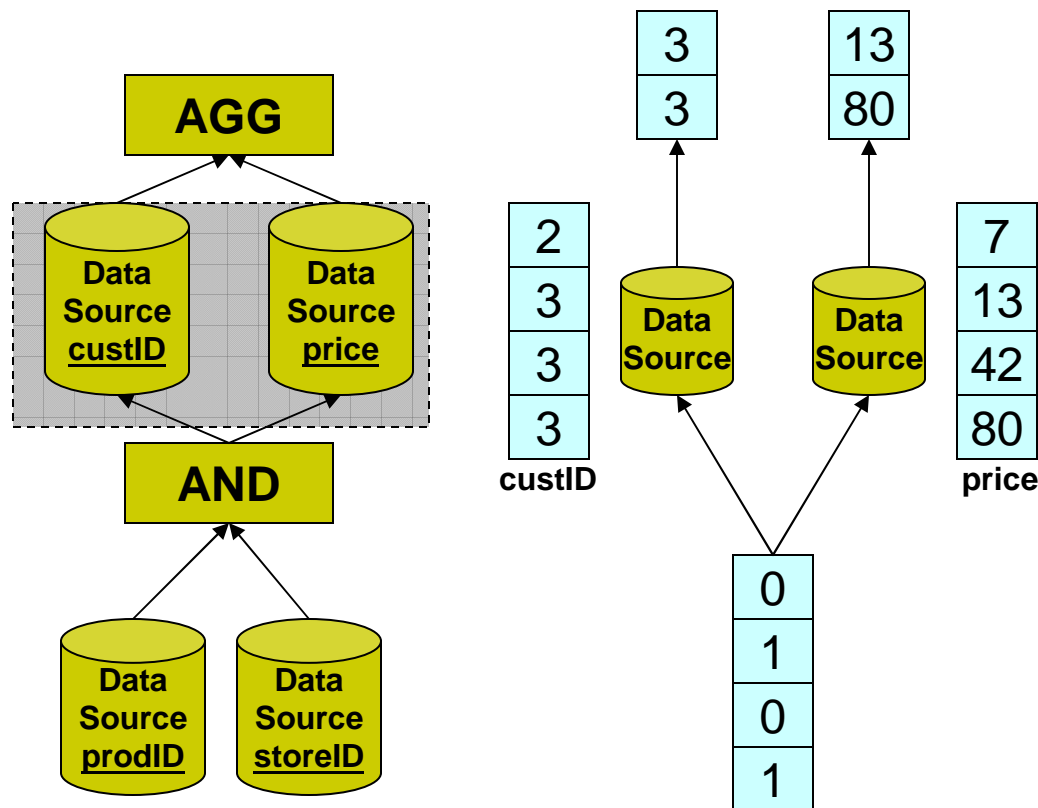
```
SELECT custID, SUM(price)
FROM table
WHERE (prodID = 4) AND
      (storeID = 1) AND
GROUP BY custID
```

prodID	storeID	custID	price
4	2	2	7
4	1	3	13
4	3	3	42
4	1	3	80





Solution 2: Operate on columns



QUERY:



```
SELECT custID, SUM(price)
FROM table
WHERE (prodID = 4) AND
      (storeID = 1) AND
GROUP BY custID
```

4	2	2	7
4	1	3	13
4	3	3	42
4	1	3	80
prodID	storeID	custID	price





Solution 2: Operate on columns



QUERY:

```
SELECT custID, SUM(price)
FROM table
WHERE (prodID = 4) AND
      (storeID = 1) AND
GROUP BY custID
```


4	2	2	7
4	1	3	13
4	3	3	42
4	1	3	80
prodID	storeID	custID	price



“Materialization Strategies in a Column-Oriented DBMS”
Abadi, Myers, DeWitt, and Madden. ICDE 2007.



For plans without joins, late materialization is a win



QUERY:

```
SELECT C1, SUM(C2)  
FROM table  
WHERE (C1 < CONST) AND  
      (C2 < CONST)  
GROUP BY C1
```


1 Ran on 2 compressed columns from TPC-H scale 10 data



“Materialization Strategies in a Column-Oriented DBMS”
Abadi, Myers, DeWitt, and Madden. ICDE 2007.



Even on uncompressed data, late materialization is still a win



QUERY:

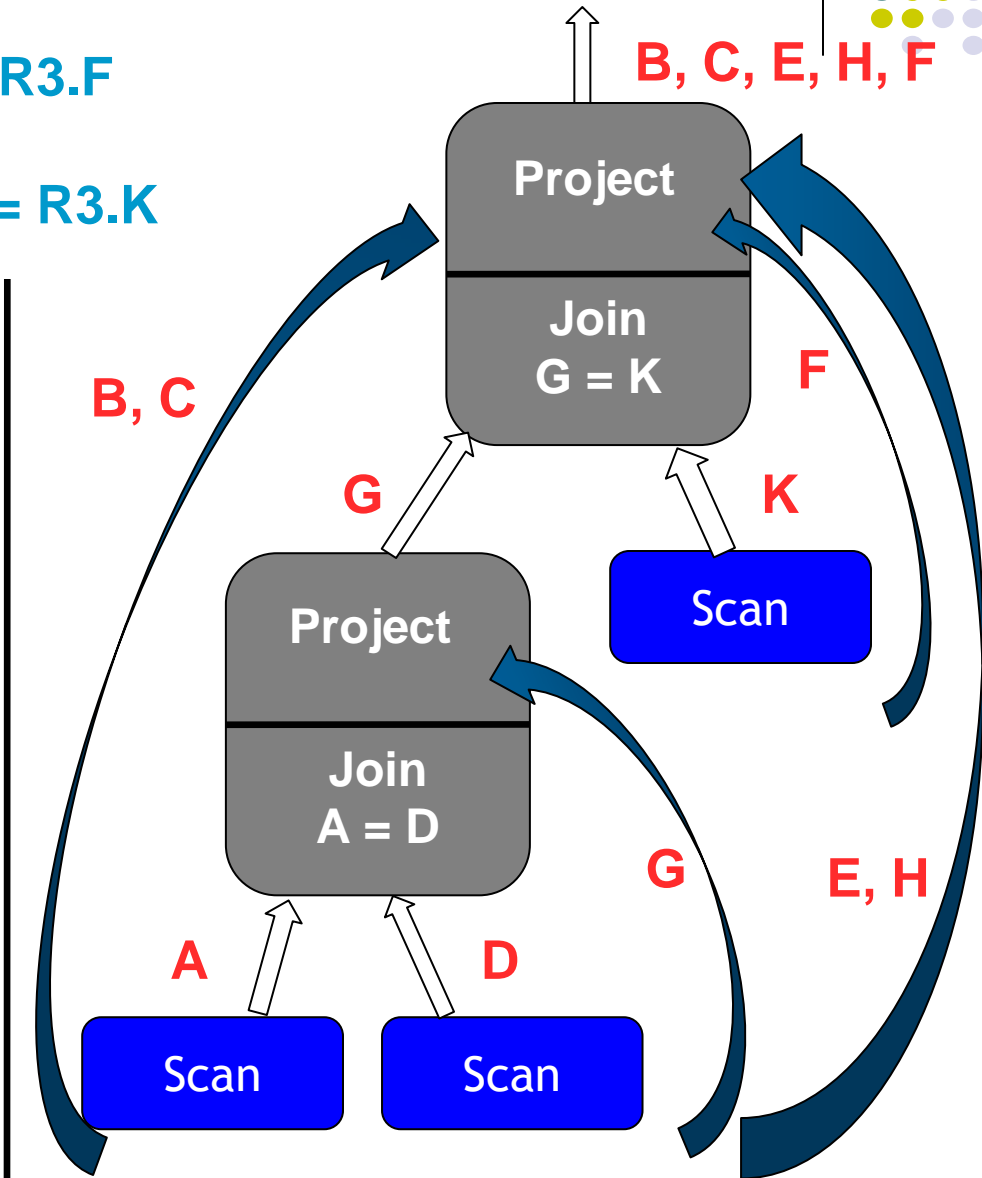

```
SELECT C1, SUM(C2)  
FROM table  
WHERE (C1 < CONST) AND  
      (C2 < CONST)  
GROUP BY C1
```

1 **Materializing late still works best**



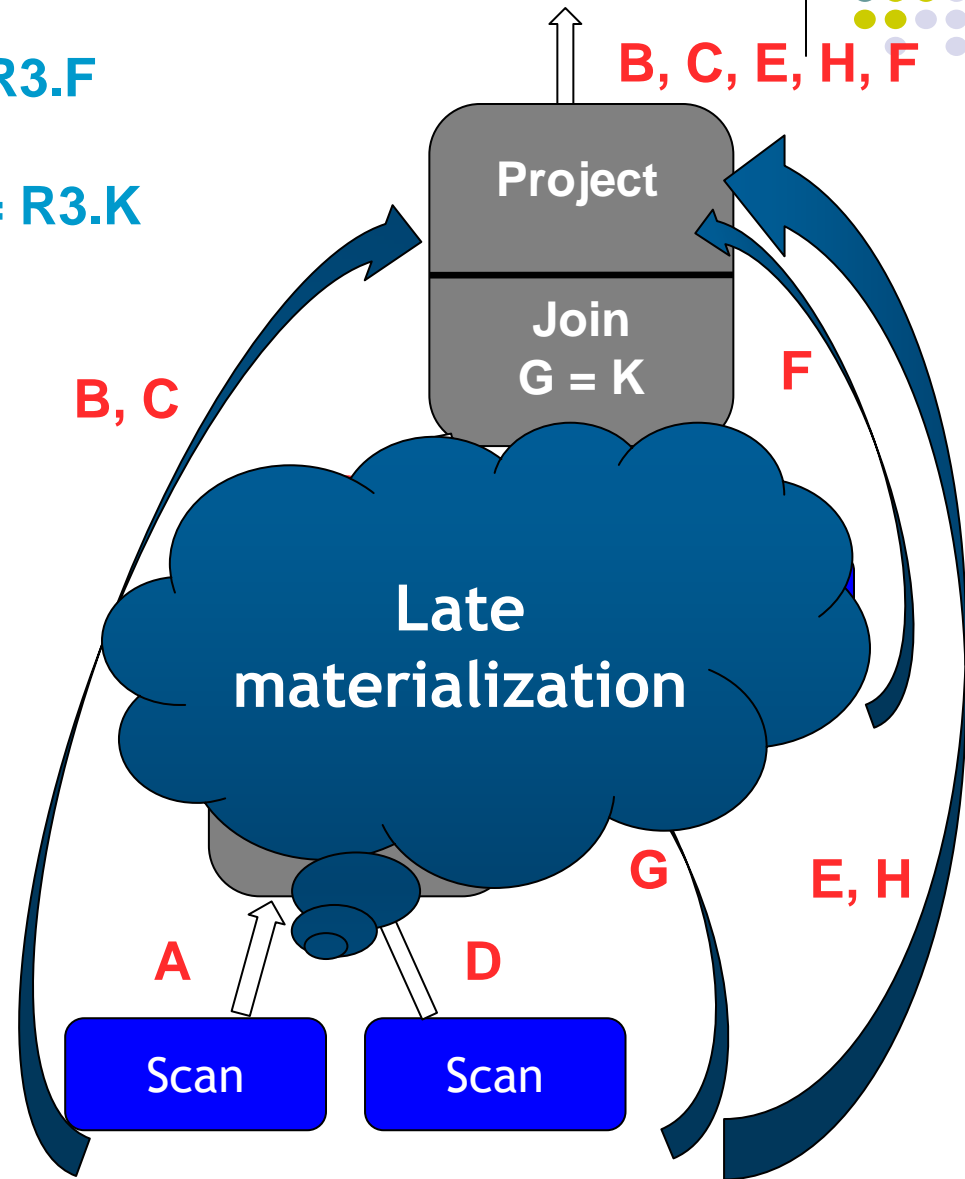

What about for plans with joins?

Select R1.B, R1.C, R2.E, R2.H, R3.F
From R1, R2, R3
Where R1.A = R2.D AND R2.G = R3.K




What about for plans with joins?

Select R1.B, R1.C, R2.E, R2.H, R3.F
From R1, R2, R3
Where R1.A = R2.D AND R2.G = R3.K






Early Materialization Example



Facts



Customers


QUERY:

```
SELECT C.lastName,SUM(F.price)
FROM facts AS F, customers AS C
WHERE F.custID = C.custID
GROUP BY C.lastName
```





Early Materialization Example



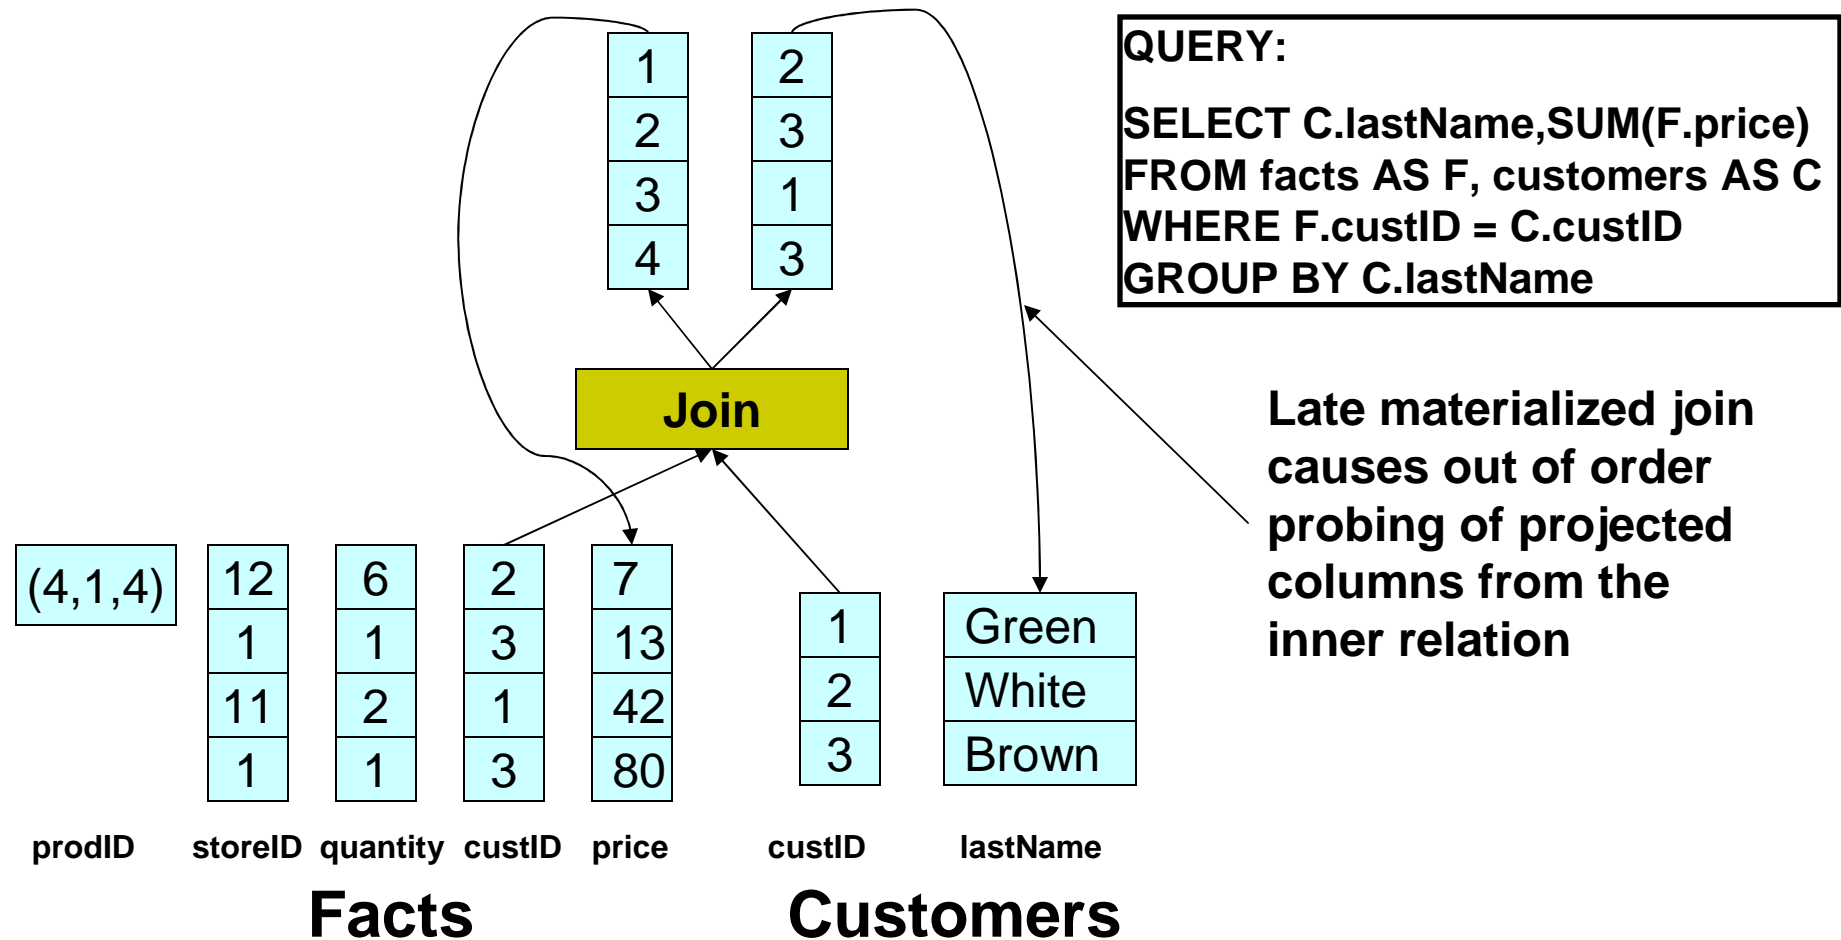
QUERY:

```
SELECT C.lastName, SUM(F.price)
FROM facts AS F, customers AS C
WHERE F.custID = C.custID
GROUP BY C.lastName
```





Late Materialization Example





Late Materialized Join Performance

- 1 Naïve LM join about 2X slower than EM join on typical queries (due to random I/O)
 - 1 This number is very dependent on
 - 1 Amount of memory available
 - 1 Number of projected attributes
 - 1 Join cardinality
- 1 But we can do better
 - 1 Invisible Join
 - 1 Jive/Flash Join
 - 1 Radix cluster/decluster join





Invisible Join

“Column-Stores vs Row-Stores: How Different are They Really?” Abadi, Madden, and Hachem. SIGMOD 2008.

- 1 Designed for typical joins when data is modeled using a star schema
 - 1 One (“fact”) table is joined with multiple dimension tables
- 1 Typical query:

```
select c_nation, s_nation, d_year,  
       sum(lo_revenue) as revenue  
from customer, lineorder, supplier, date  
where lo_custkey = c_custkey  
      and lo_suppkey = s_suppkey  
      and lo_orderdate = d_datekey  
      and c_region = 'ASIA'  
      and s_region = 'ASIA'  
      and d_year >= 1992 and d_year <= 1997  
group by c_nation, s_nation, d_year  
order by d_year asc, revenue desc;
```





Invisible Join

“Column-Stores vs Row-Stores: How Different are They Really?” Abadi, Madden, and Hachem. SIGMOD 2008.

Apply “region = ‘Asia’” On Customer Table

custkey	region	nation	...
1	ASIA	CHINA	...
2	ASIA	INDIA	...
3	ASIA	INDIA	...
4	EUROPE	FRANCE	...



Hash Table (or bit-map)
Containing Keys 1, 2 and 3

Apply “region = ‘Asia’” On Supplier Table

suppkey	region	nation	...
1	ASIA	RUSSIA	...
2	EUROPE	SPAIN	...
3	ASIA	JAPAN	...



Hash Table (or bit-map)
Containing Keys 1, 3

Apply “year in [1992,1997]” On Date Table

dateid	year	...
01011997	1997	...
01021997	1997	...
01031997	1997	...




Hash Table Containing
Keys 01011997, 01021997,
and 01031997



Original Fact Table

orderkey	custkey	suppkey	orderdate	revenue
1	3	1	01011997	43256
2	3	2	01011997	33333
3	4	3	01021997	12121
4	1	1	01021997	23233
5	4	2	01021997	45456
6	1	2	01031997	43251
7	3	2	01031997	34235

“Column-Stores vs Row-Stores:
How Different are They Really?”
Abadi et. al. SIGMOD 2008.



Hash Table
Containing
Keys 1, 2 and 3

+

custkey
3
3
4
1
4
1
3

Hash Table
Containing
Keys 1 and 3

+

suppkey
1
2
3
1
2
2
2

Hash Table
Containing
Keys 01011997,
01021997, and 01031997

+




orderdate
01011997
01011997
01021997
01021997
01021997
01031997
01031997

1
1
0
1
0
1
1

1
0
1
1
0
0
0

1
1
1
1
1
1
1





“Column-Stores vs Row-Stores: How Different are They Really?” Abadi, Madden, and Hachem. SIGMOD 2008.



Invisible Join

Apply “region = ‘Asia’” On Customer Table

custkey	region	nation	...
1	ASIA	CHINA	...
2	ASIA	INDIA	...
3	ASIA	INDIA	...
4	EUROPE	FRANCE	...



~~Hash Table (or bit-map)
Containing Keys 1, 2 and 3~~

Range [1-3]
(between-predicate rewriting)

Apply “region = ‘Asia’” On Supplier Table

suppkey	region	nation	...
1	ASIA	RUSSIA	...
2	EUROPE	SPAIN	...
3	ASIA	JAPAN	...



Hash Table (or bit-map)
Containing Keys 1, 3

Apply “year in [1992,1997]” On Date Table

dateid	year	...
01011997	1997	...
01021997	1997	...
01031997	1997	...



Hash Table Containing
Keys 01011997, 01021997,
and 01031997






Invisible Join



1 Bottom Line

- 1 Many data warehouses model data using star/snowflake schemes
- 1 Joins of one (fact) table with many dimension tables is common
- 1 Invisible join takes advantage of this by making sure that the table that can be accessed in position order is the fact table for each join
- 1 Position lists from the fact table are then intersected (in position order)
- 1 This reduces the amount of data that must be accessed out of order from the dimension tables
- 1 “Between-predicate rewriting” trick not relevant for this discussion



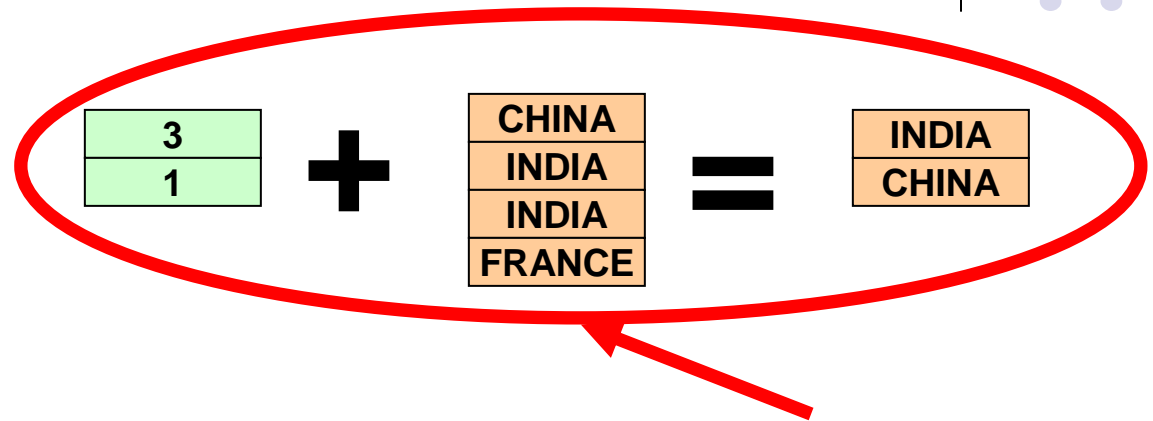


Still accessing table out of order





Jive/Flash Join




Still accessing table out of order

“Fast Joins using Join Indices”. Li and Ross, VLDBJ 8:1-24, 1999.

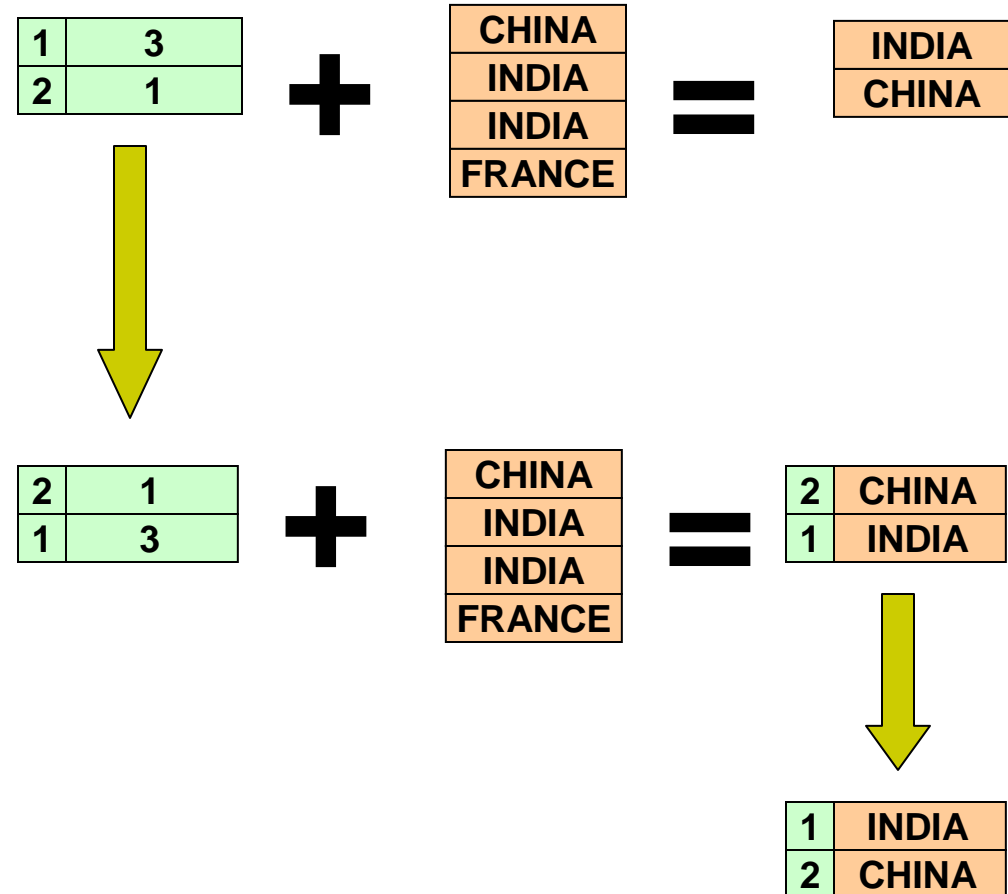
“Query Processing Techniques for Solid State Drives”. Tsirogiannis, Harizopoulos et. al. SIGMOD 2009.





Jive/Flash Join

1. Add column with dense ascending integers from 1
2. Sort new position list by second column
3. Probe projected column in order using new sorted position list, keeping first column from position list around
4. Sort new result by first column





Jive/Flash Join

1 Bottom Line

- 1 Instead of probing projected columns from inner table out of order:
 - 1 Sort join index
 - 1 Probe projected columns in order
 - 1 Sort result using an added column
- 1 LM vs EM tradeoffs:
 - 1 LM has the extra sorts (EM accesses all columns in order)
 - 1 LM only has to fit join columns into memory (EM needs join columns and all projected columns)
 - § Results in big memory and CPU savings (see part 3 for why there is CPU savings)
 - 1 LM only has to materialize relevant columns
 - 1 In many cases LM advantages outweigh disadvantages
- 1 LM would be a clear winner if not for those pesky sorts ... can we do better?





Radix Cluster/Decluster

- 1 The full sort from the Jive join is actually overkill
 - 1 We just want to access the storage blocks in order (we don't mind random access within a block)
 - 1 So do a radix sort and stop early
 - 1 By stopping early, data within each block is accessed out of order, but in the order specified in the original join index
 - 1 Use this pseudo-order to accelerate the post-probe sort as well

•“Database Architecture Optimized for the New Bottleneck: Memory Access”
VLDB'99
•“Generic Database Cost Models for Hierarchical Memory Systems”, VLDB'02
(all Manegold, Boncz, Kersten)

“Cache-Conscious Radix-Decluster Projections”, Manegold, Boncz, Nes,
VLDB'04





Radix Cluster/Deccluster

- 1 Bottom line
 - 1 Both sorts from the Jive join can be significantly reduced in overhead
 - 1 Only been tested when there is sufficient memory for the entire join index to be stored three times
 - 1 Technique is likely applicable to larger join indexes, but utility will go down a little
 - 1 Only works if random access within a storage block
 - 1 Don't want to use radix cluster/decluster if you have variable-width column values or compression schemes that can only be decompressed starting from the beginning of the block





LM vs EM joins

- 1 Invisible, Jive, Flash, Cluster, Decluster techniques contain a bag of tricks to improve LM joins
- 1 Research papers show that LM joins become 2X faster than EM joins (instead of 2X slower) for a wide array of query types





Tuple Construction Heuristics

- 1 For queries with selective predicates, aggregations, or compressed data, use late materialization
- 1 For joins:
 - 1 Research papers:
 - 1 Always use late materialization
 - 1 Commercial systems:
 - 1 Inner table to a join often materialized before join (reduces system complexity):
 - 1 Some systems will use LM only if columns from inner table can fit entirely in memory





Outline

- 1 Computational Efficiency of DB on modern hardware
 - 1 how column-stores can help here
 - 1 Keynote revisited: MonetDB & VectorWise in more depth
- 1 CPU efficient column compression
 - 1 vectorized decompression
- 1 Conclusions
 - 1 future work




Column-Oriented Database Systems

VLDB
2009
Tutorial



40 years of hardware evolution
vs.
DBMS computational efficiency






CPU Architecture

Elements:


- 1 Storage
 - 1 CPU caches L1/L2/L3
- 1 Registers
- 1 Execution Unit(s)
 - 1 Pipelined
 - 1 SIMD





CPU Metrics

Processor	16-bit address/ bus, micro-coded	32-bit address/ bus, micro-coded
Product	80286	80386
Year	1982	1985
Transistors (thousands)	134	275
Latency (clocks)	6	10
Bus width (bits)	16	32
Clock rate (MHz)	12.5	10
Bandwidth (MIPS)	2	3
Latency (ns)	320	300



core
Duo
6
300
3
00






CPU Metrics

Processor	16-bit address/ bus, micro-coded	32-bit address/ bus, micro-coded	5-stage pipeline, on-chip I&D caches FPU	2-way super-scalar, 64-bit bus	Out-of-order, 3-way super-scalar	Out-of-order, super-pipelined, on-chip L2 cache	Multi-core
Product	80286	80386	80486	Pentium	PentiumPro	Pentium4	CoreDuo
Year	1982	1985	1989	1993	1997	2001	2006
Transistors (thousands)	134	275	1,200	3,100	5,500	42,000	151,600
Latency (clocks)	6	5	5	5	10	22	12
Bus width (bits)	16	32	32	64	64	64	64
Clock rate (MHz)	12.5	16	25	66	200	1500	2333
Bandwidth (MIPS)	2	6	25	132	600	4500	21000
Latency (ns)	320	313	200	76	50	15	5






DRAM Metrics





Super-Scalar Execution (pipelining)





Hazards

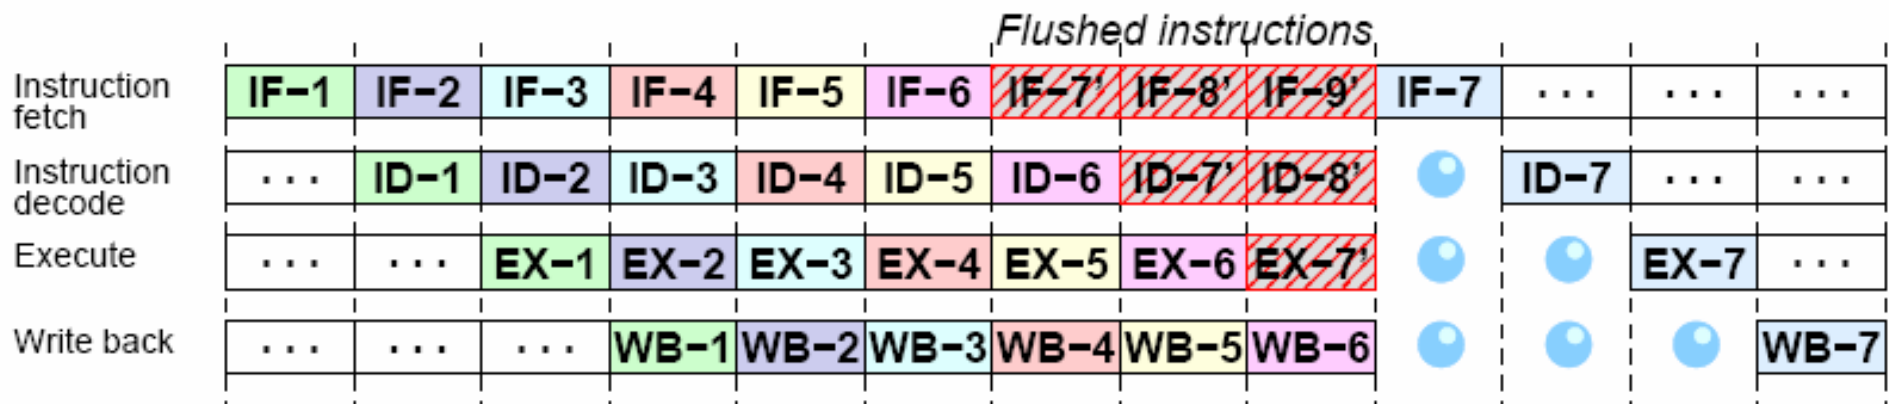
1 Data hazards

- 1 Dependencies between instructions
- 1 L1 data cache misses

1 Control Hazards

- 1 Branch mispredictions
- 1 Computed branches (late binding)
- 1 L1 instruction cache misses

Result: bubbles in the pipeline




Out-of-order execution addresses data hazards

- 1 control hazards typically more expensive





SIMD




- 1 Single Instruction Multiple Data
 - 1 Same operation applied on a vector of values
 - 1 MMX: 64 bits, SSE: 128bits, AVX: 256bits
 - 1 SSE, e.g. multiply 8 short integers





A Look at the Query Pipeline




```
SELECT id, name
      (age-30)*50 AS bonus
FROM   employee
WHERE  age > 30
```





A Look at the Query Pipeline



Operators

Iterator interface

-open()


-**next()**: tuple

-close()





A Look at the Query Pipeline



Primitives

Provide computational functionality

All arithmetic allowed in expressions, e.g. Multiplication

$7 * 50$


`mult(int,int) ↦ int`





Database Architecture causes Hazards

- DB workload execution on a modern computer



“DBMSs On A Modern Processor: Where Does Time Go? ”
Ailamaki, DeWitt, Hill, Wood, VLDB’99





DBMS Computational Efficiency

TPC-H 1GB, query 1

1 selects 98% of fact table, computes net prices and aggregates all

1 Results:

1 C program:	?
1 MySQL:	26.2s
1 DBMS “X”:	28.1s

“MonetDB/X100: Hyper-Pipelining Query Execution ” Boncz, Zukowski, Nes, CIDR’05





DBMS Computational Efficiency

TPC-H 1GB, query 1

1 selects 98% of fact table, computes net prices and aggregates all

1 Results:


1 C program:	0.2s
1 MySQL:	26.2s
1 DBMS “X”:	28.1s

“MonetDB/X100: Hyper-Pipelining Query Execution ” Boncz, Zukowski, Nes, CIDR’05



Column-Oriented Database Systems

VLDB
2009
Tutorial






a column-store

- 1 “save disk I/O when scan-intensive queries need a few columns”
- 1 “avoid an expression interpreter to improve computational efficiency”





RISC Database Algebra



```

SELECT  id, name, (age-30)*50 as bonus
FROM    people
WHERE   age > 30
    
```





RISC Database Algebra

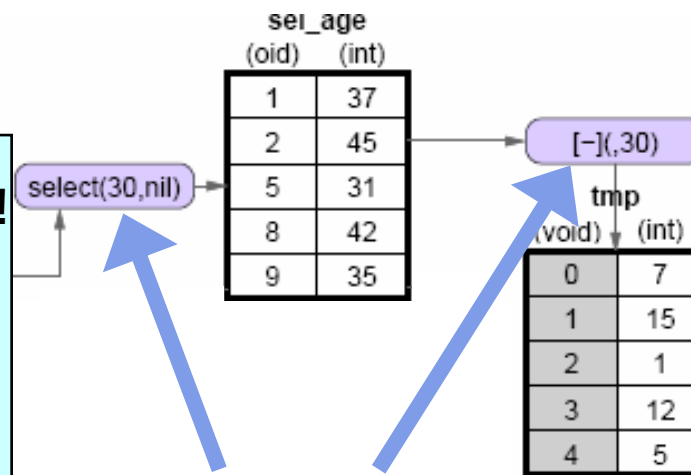
CPU happy? Give it “nice” code !

- few dependencies (control,data)
- CPU gets out-of-order execution
- compiler can e.g. generate SIMD

One loop for an entire column

- no per-tuple interpretation
- arrays: no record navigation
- better instruction cache locality

```
{  
    for(i=0; i<n; i++)  
        res[i] = col[i] - val;  
}
```



Simple, hard-coded semantics in operators





RISC Database Algebra

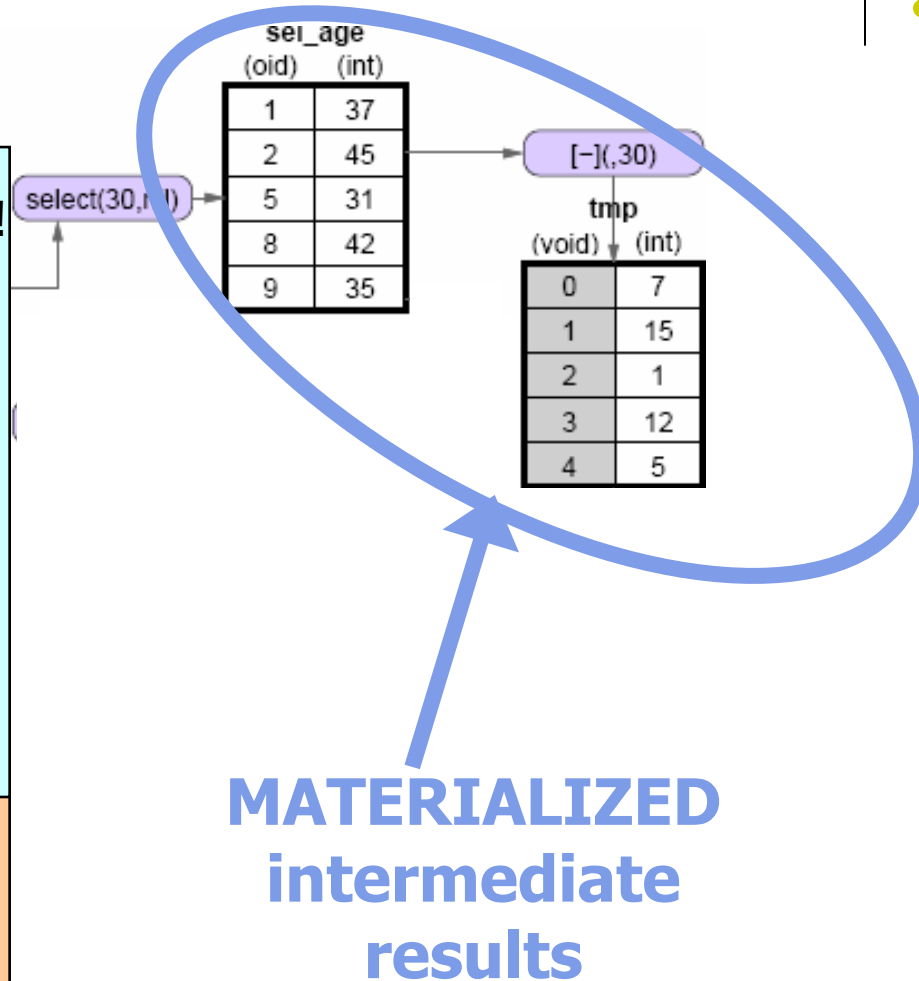
CPU happy? Give it “nice” code !

- few dependencies (control,data)
- CPU gets out-of-order execution
- compiler can e.g. generate SIMD

One loop for an entire column

- no per-tuple interpretation
- arrays: no record navigation
- better instruction cache locality

```
{  
    for(i=0; i<n; i++)  
        res[i] = col[i] - val;  
}
```





a column-store



- 1 “save disk I/O when scan-intensive queries need a few columns”
- 1 “avoid an expression interpreter to improve computational efficiency”

How?

- 1 RISC query algebra: hard-coded semantics
 - 1 Decompose complex expressions in multiple operations
- 1 Operators only handle **simple arrays**
 - 1 No code that handles slotted buffered record layout
- 1 Relational algebra becomes **array manipulation language**
 - 1 Often SIMD for free
- 1 Plus: use of *cache-conscious* algorithms for Sort/Aggr/Join





a Faustian pact

- 1 You want efficiency
 - 1 Simple hard-coded operators
- 1 I take scalability
 - 1 Result materialization

n	C program:	0.2s
n	MonetDB:	3.7s
n	MySQL:	26.2s
n	DBMS "X":	28.1s




Column-Oriented Database Systems

VLDB
2009
Tutorial



as a research platform





SIGMOD 1985

A DECOMPOSITION STORAGE MODEL

George P. Copeland
Setrag N.

MonetDB
BAT Algebra

MonetDB supports
SQL, XML, ODMG,
...RDF

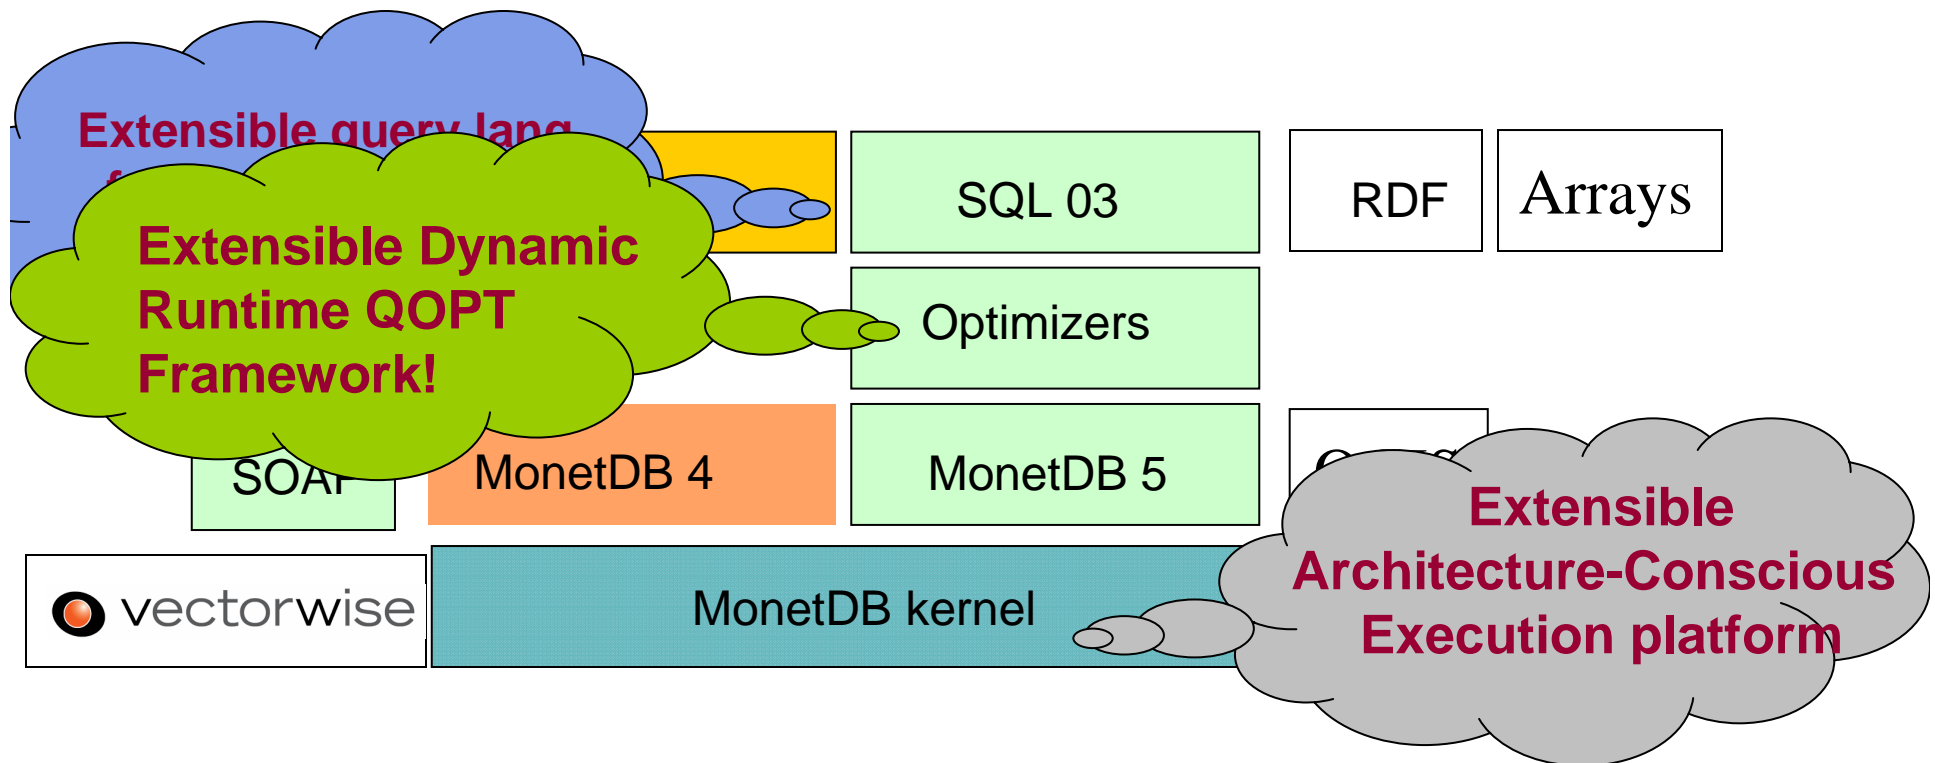
RDF support on C-
STORE / SW-Store

- “MIL Primitives for Querying a Fragmented World”, Boncz, Kersten, VLDBJ’98
- “Flattening an Object Algebra to Provide Performance” Boncz, Wilschut, Kersten, ICDE’98
- “MonetDB/XQuery: a fast XQuery processor powered by a relational engine” Boncz, Grust, vanKeulen, Rittinger, Teubner, SIGMOD’06
- “SW-Store: a vertically partitioned DBMS for Semantic Web data management” Abadi, Marcus, Madden, Hollenbach, VLDBJ’09





The MONETDB Software Stack





optimizer

frontend

backend

as a research platform



1 Cache-Conscious Joins

- 1 Cost Models, Radix-cluster
- Radix-decluster

- “Database Architecture Optimized for the New Bottleneck: Memory Access” VLDB’99
- “Generic Database Cost Models for Hierarchical Memory Systems”, VLDB’02 (all Manegold, Boncz, Kersten)
- “Cache-Conscious Radix-Decluster Projections”, Manegold, Boncz, Nes, VLDB’04

1 MonetDB/XQuery:

- 1 structural joins exploiting positional column access

“MonetDB/XQuery: a fast XQuery processor powered by a relational engine” Boncz, Grust, vanKeulen, Rittinger, Teubner, SIGMOD’06

1 Cracking:

- 1 on-the-fly automatic indexing without workload knowledge

“Database Cracking”, CIDR’07
 “Updating a cracked database “, SIGMOD’07
 “Self-organizing tuple reconstruction in column-stores“, SIGMOD’09 (all Idreos, Manegold, Kersten)

1 Recycling:

- 1 using materialized intermediates

“An architecture for recycling intermediates in a column-store”, Ivanova, Kersten, Nes, Goncalves, SIGMOD’09

1 Run-time Query Optimization:

- 1 correlation-aware run-time optimization without cost model

“ROX: run-time optimization of XQueries”, Abdelkader, Boncz, Manegold, vanKeulen, SIGMOD’09



Column-Oriented Database Systems


VLDB
2009
Tutorial



 vectorwise

“MonetDB/X100”

vectorized query processing





The diagram illustrates the execution of a query in MonetDB, highlighting materialized intermediate results. A red oval encircles the join operation and the temporary table 'tmp', with an arrow pointing to the text 'MATERIALIZED intermediate results'.

Query Plan Components:

- select(30, ...)**: A selection node that filters rows based on a condition involving the value 30.
- set_age**: A table with columns 'oid' and 'age'.

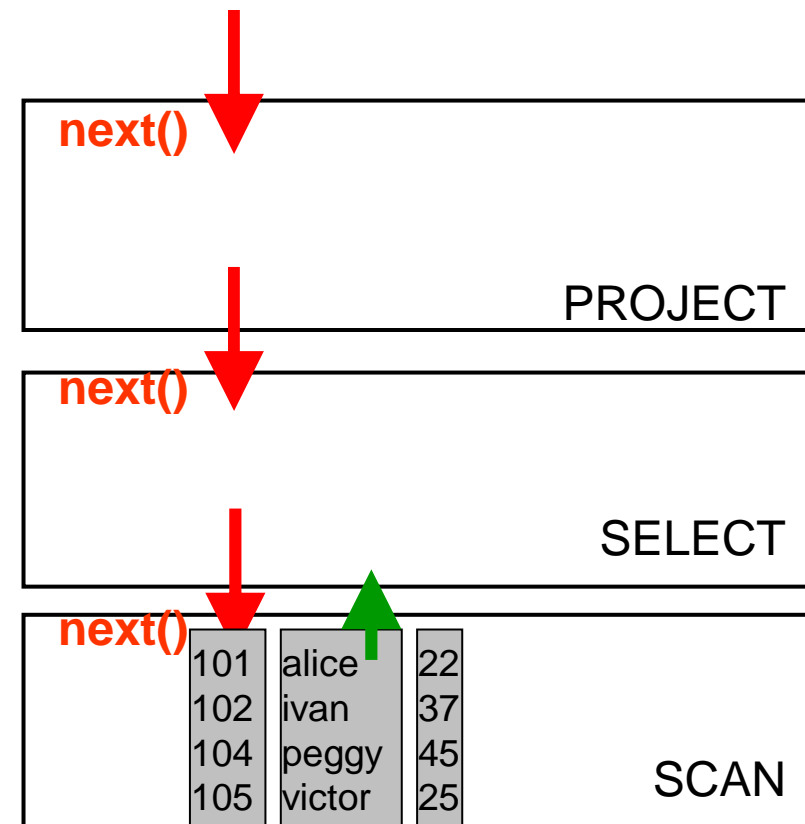
oid	age
1	37
2	45
5	31
8	42
9	35
- Join Operation**: Represented by the node **(-),(30)**, indicating a join with a constant value of 30.
- tmp**: A temporary table that stores the result of the join operation.

(void)	(int)
0	7
1	15
2	1
3	12
4	5

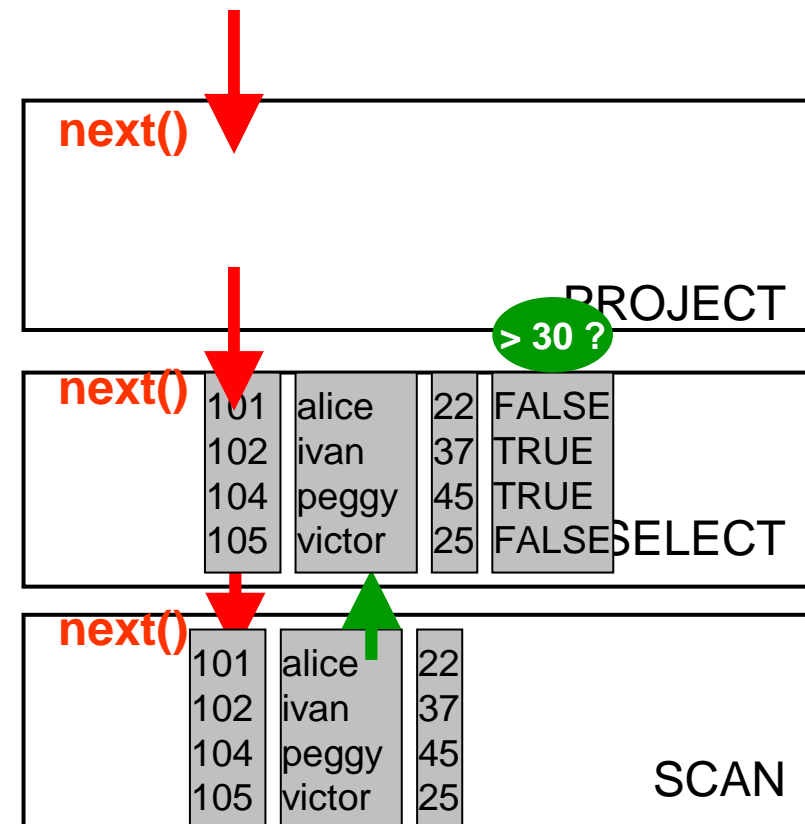
MATERIALIZED intermediate results



“MonetDB/X100: Hyper-Pipelining Query Execution” Boncz, Zukowski, Nes, CIDR’05



“MonetDB/X100: Hyper-Pipelining Query Execution” Boncz, Zukowski, Nes, CIDR’05



“MonetDB/X100: Hyper-Pipelining Query Execution” Boncz, Zukowski, Nes, CIDR’05

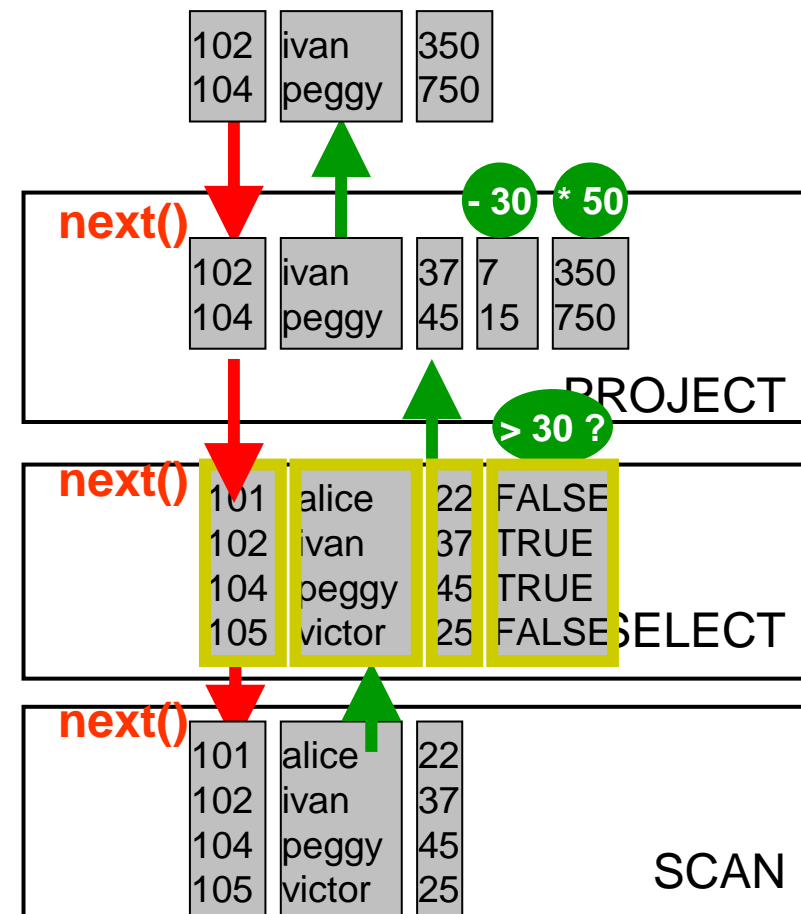


“Vectorized In Cache Processing”

vector = array of ~100

processed in a tight loop

CPU cache Resident



“MonetDB/X100: Hyper-Pipelining Query Execution” Boncz, Zukowski, Nes, CIDR’05



Observations:

`next()` called much less often \perp
 more time spent in **primitives**
 less in **overhead**

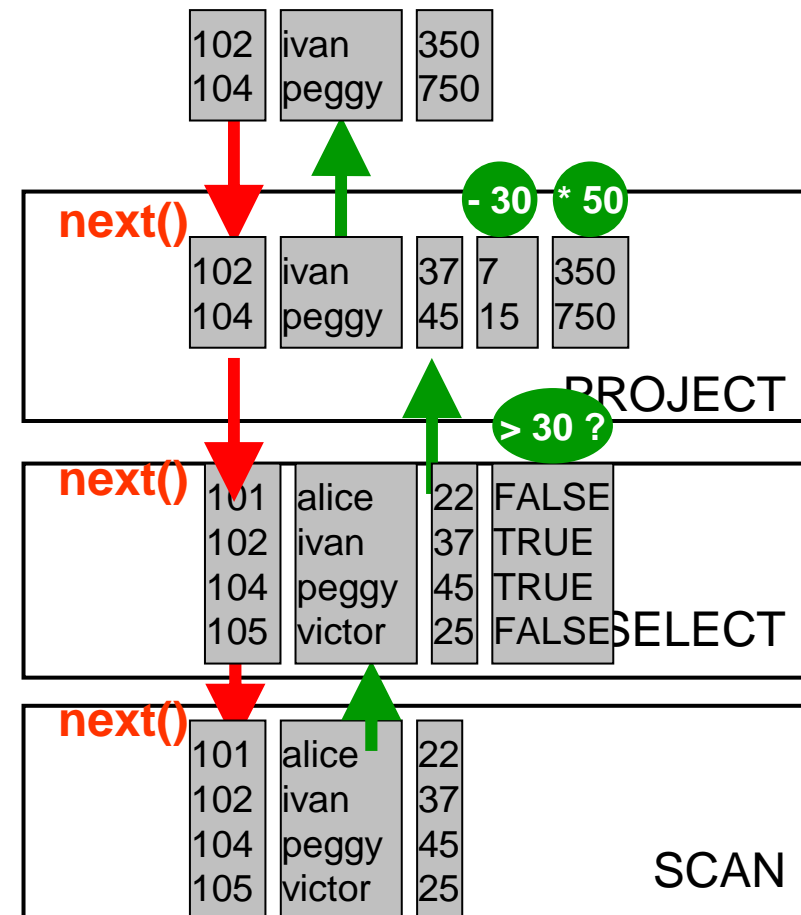
primitive calls process an array of values in a **loop**:

CPU Efficiency depends on “nice” code

- out-of-order execution
- few dependencies (control,data)
- compiler support

Compilers like simple loops over arrays

- loop-pipelining
- automatic SIMD



“MonetDB/X100: Hyper-Pipelining Query Execution” Boncz, Zukowski, Nes, CIDR’05



Observations:

`next()` called much less often \perp
more time spent in **primitives**
less in **overhead**

primitive calls process an array of values in a **loop**:

CPU Efficiency depends on “nice” code

- out-of-order execution
- few dependencies (control,data)
- compiler support

Compilers like simple loops over arrays

- loop-pipelining
- automatic SIMD

> 30 ?

FALSE
TRUE
TRUE
FALSE

```
for(i=0; i<n; i++)  
    res[i] = (col[i] > x)
```

- 30

7
15

```
for(i=0; i<n; i++)  
    res[i] = (col[i] - x)
```

* 50

350
750

```
for(i=0; i<n; i++)  
    res[i] = (col[i] * x)
```




“MonetDB/X100: Hyper-Pipelining Query Execution ” Boncz, Zukowski, Nes, CIDR’05



Tricks being played:

- Late materialization
- Materialization avoidance using **selection vectors**




“MonetDB/X100: Hyper-Pipelining Query Execution” Boncz, Zukowski, Nes, CIDR’05



```
map_mulflt_valflt_col(
    float *res,
    int* sel,
    float val,
    float *col, int n)
{
    for(int i=0; i<n; i++)
        res[i] = val * col[sel[i]];
}
```

selection vectors used to reduce vector copying

contain selected positions




“MonetDB/X100: Hyper-Pipelining Query Execution” Boncz, Zukowski, Nes, CIDR’05



```
map_mulflt_valflt_col(
    float *res,
    int* sel,
    float val,
    float *col, int n)
{
    for(int i=0; i<n; i++)
        res[i] = val * col[sel[i]];
}
```

selection vectors used to reduce vector copying

contain selected positions





MonetDB/X100

“MonetDB/X100: Hyper-Pipelining Query Execution” Boncz, Zukowski, Nes, CIDR’05



- 1 Both efficiency
 - 1 Vectorized primitives
- 1 and scalability..
 - 1 Pipelined query evaluation

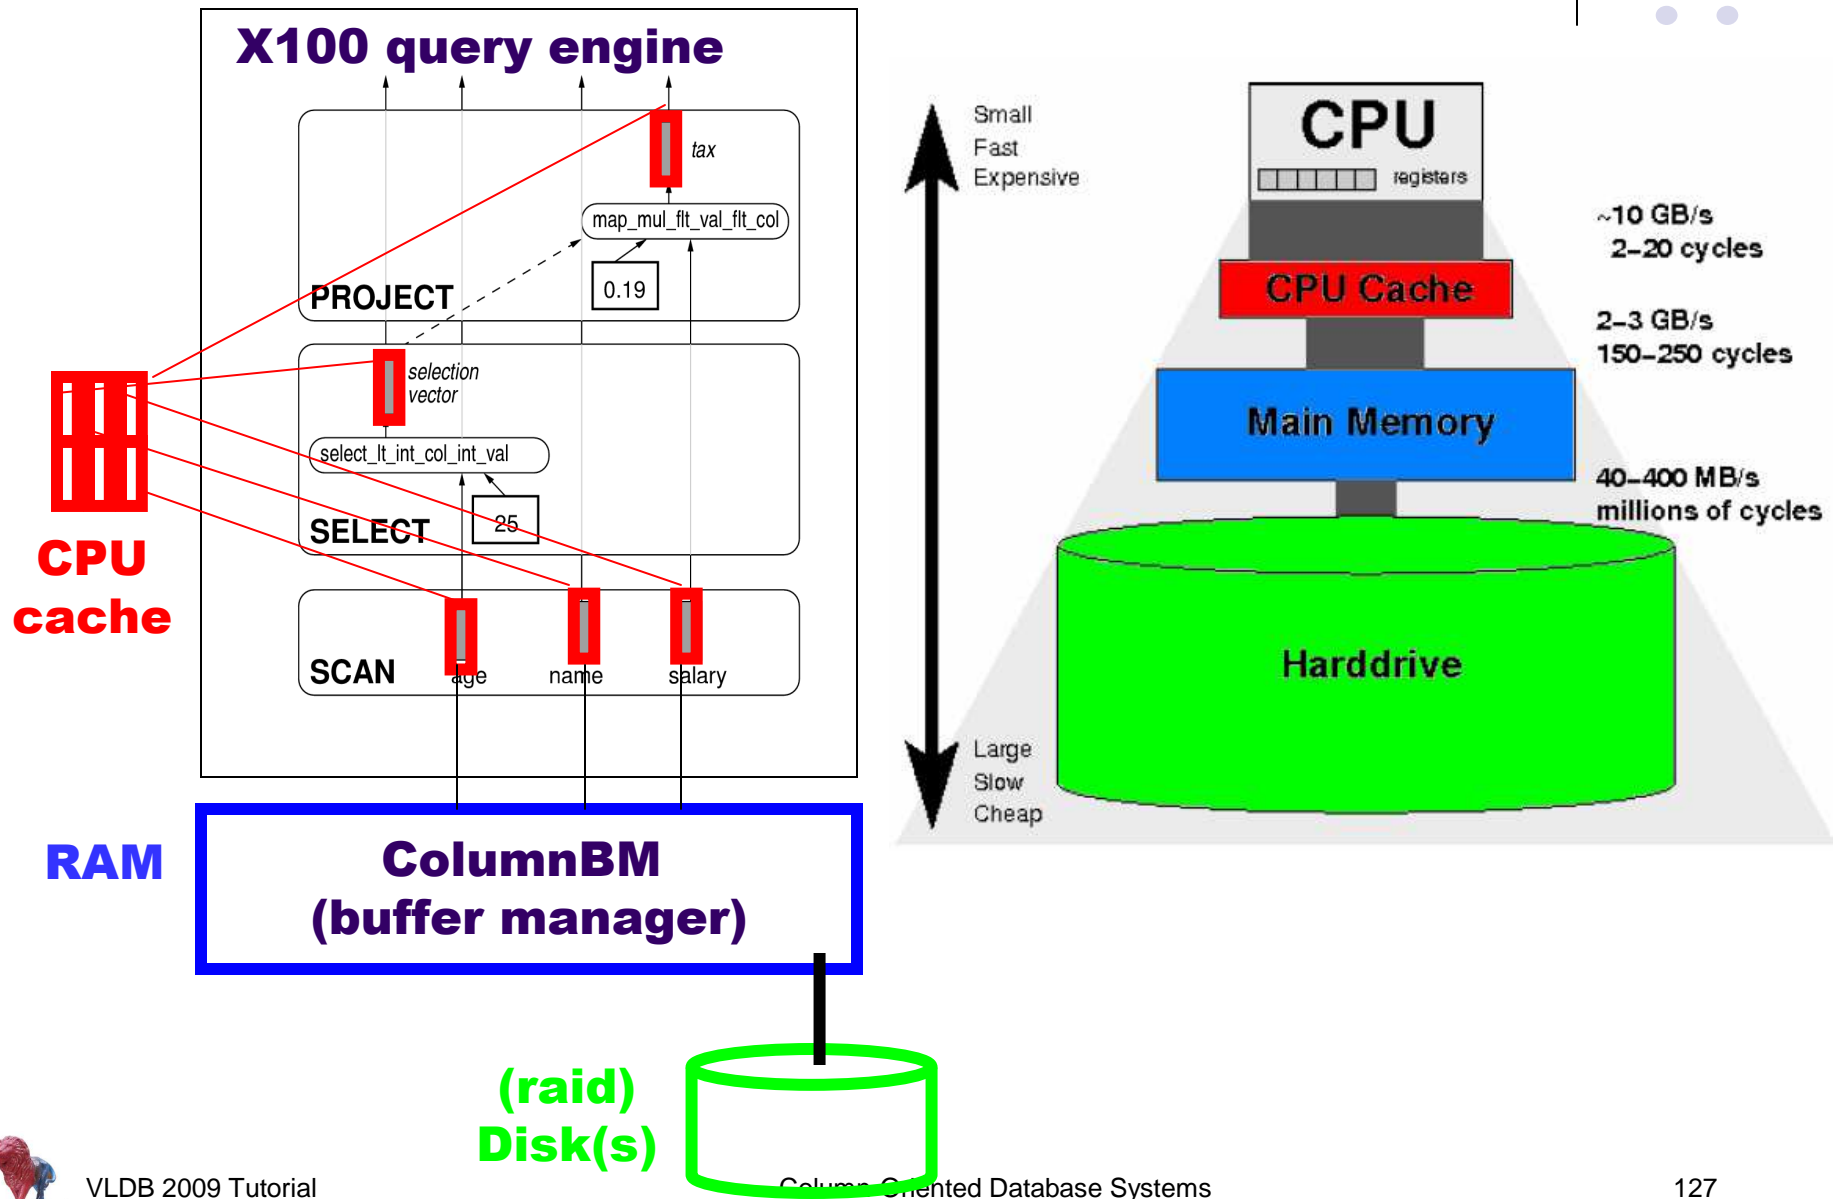
n C program:	0.2s
n MonetDB/X100:	0.6s
n MonetDB:	3.7s
n MySQL:	26.2s
n DBMS “X”:	28.1s





Memory Hierarchy

“MonetDB/X100: Hyper-Pipelining Query Execution” Boncz, Zukowski, Nes, CIDR’05

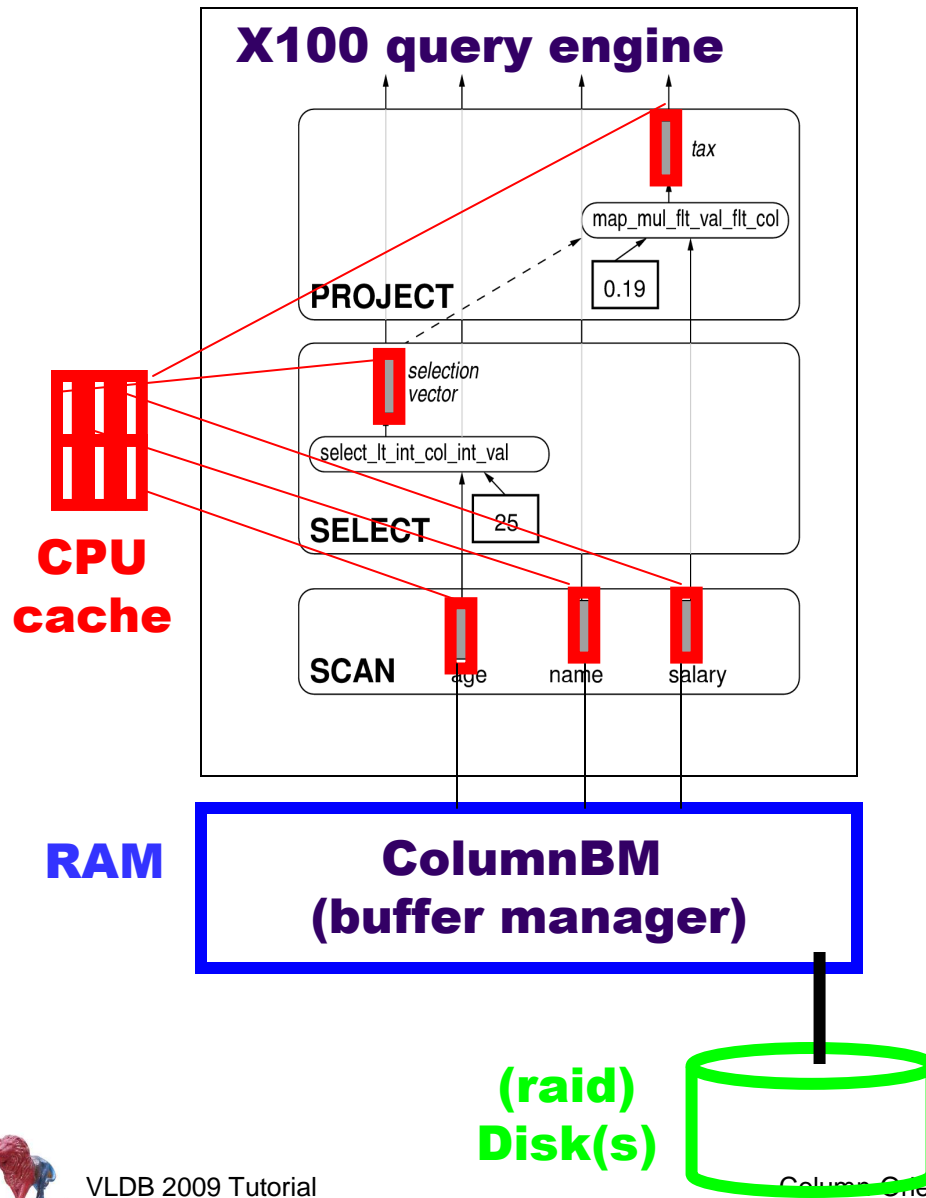




“MonetDB/X100: Hyper-Pipelining Query Execution” Boncz, Zukowski, Nes, CIDR’05



Memory Hierarchy



Vectors are only the in-cache representation

RAM & disk representation might actually be different

(vectorwise uses both PAX & DSM)






“MonetDB/X100: Hyper-Pipelining Query Execution” Boncz, Zukowski, Nes, CIDR’05



Optimal Vector size?




All vectors together should fit the CPU cache

Optimizer should tune this, given the query characteristics.






Varying the Vector size



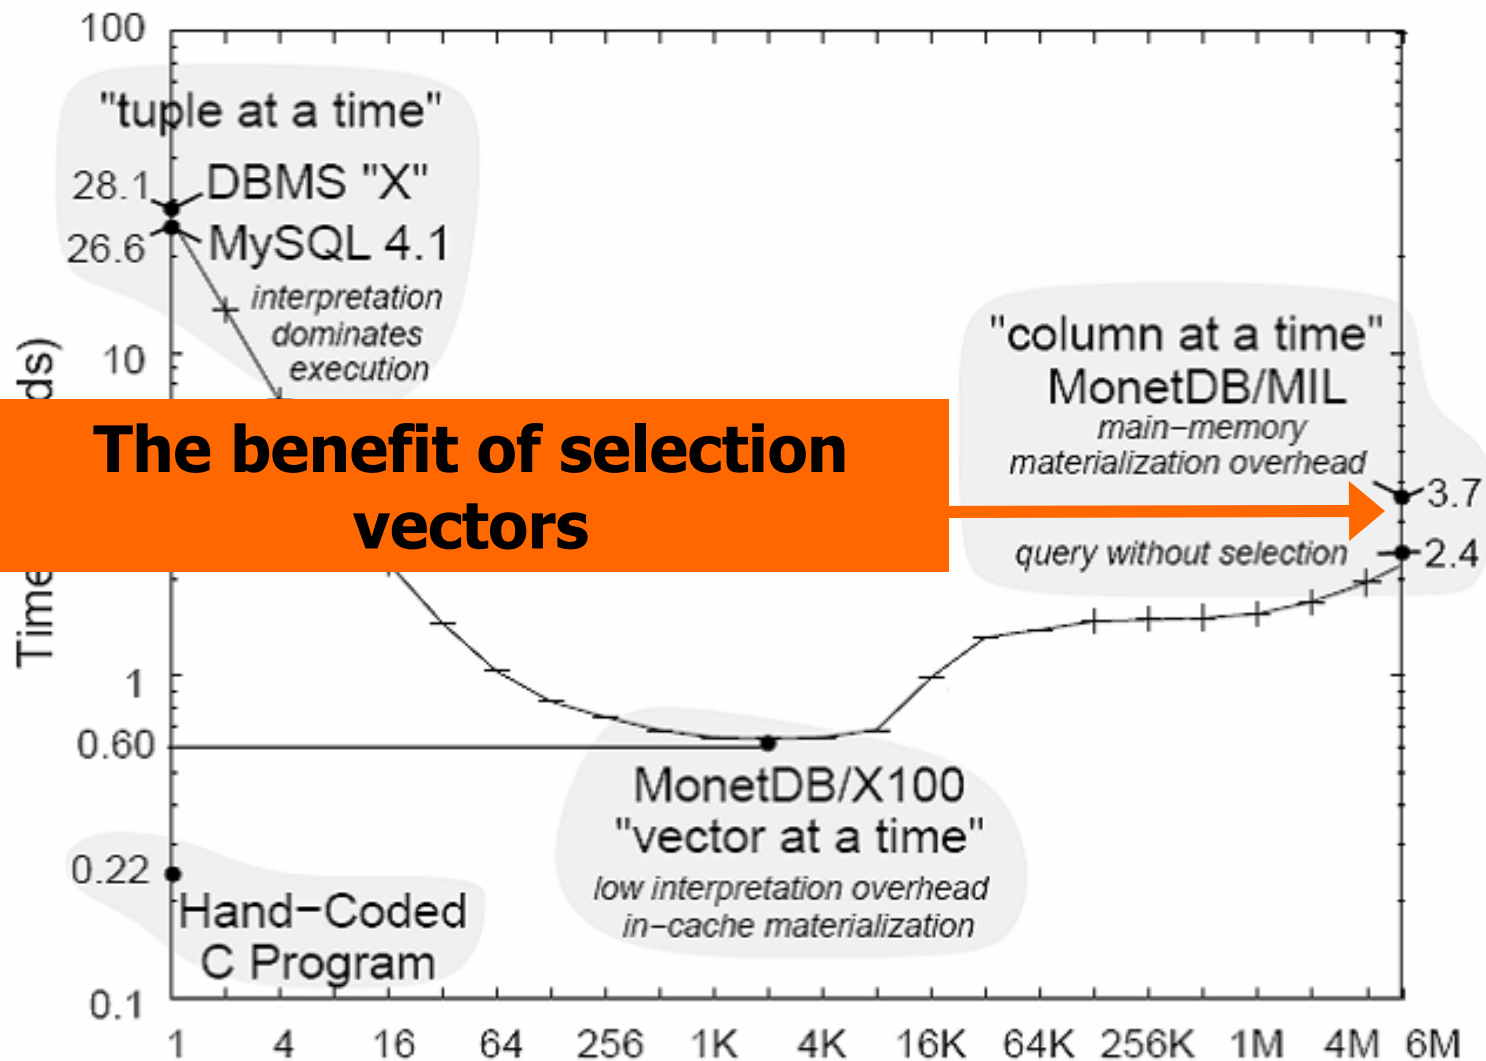


Varying the Vector size



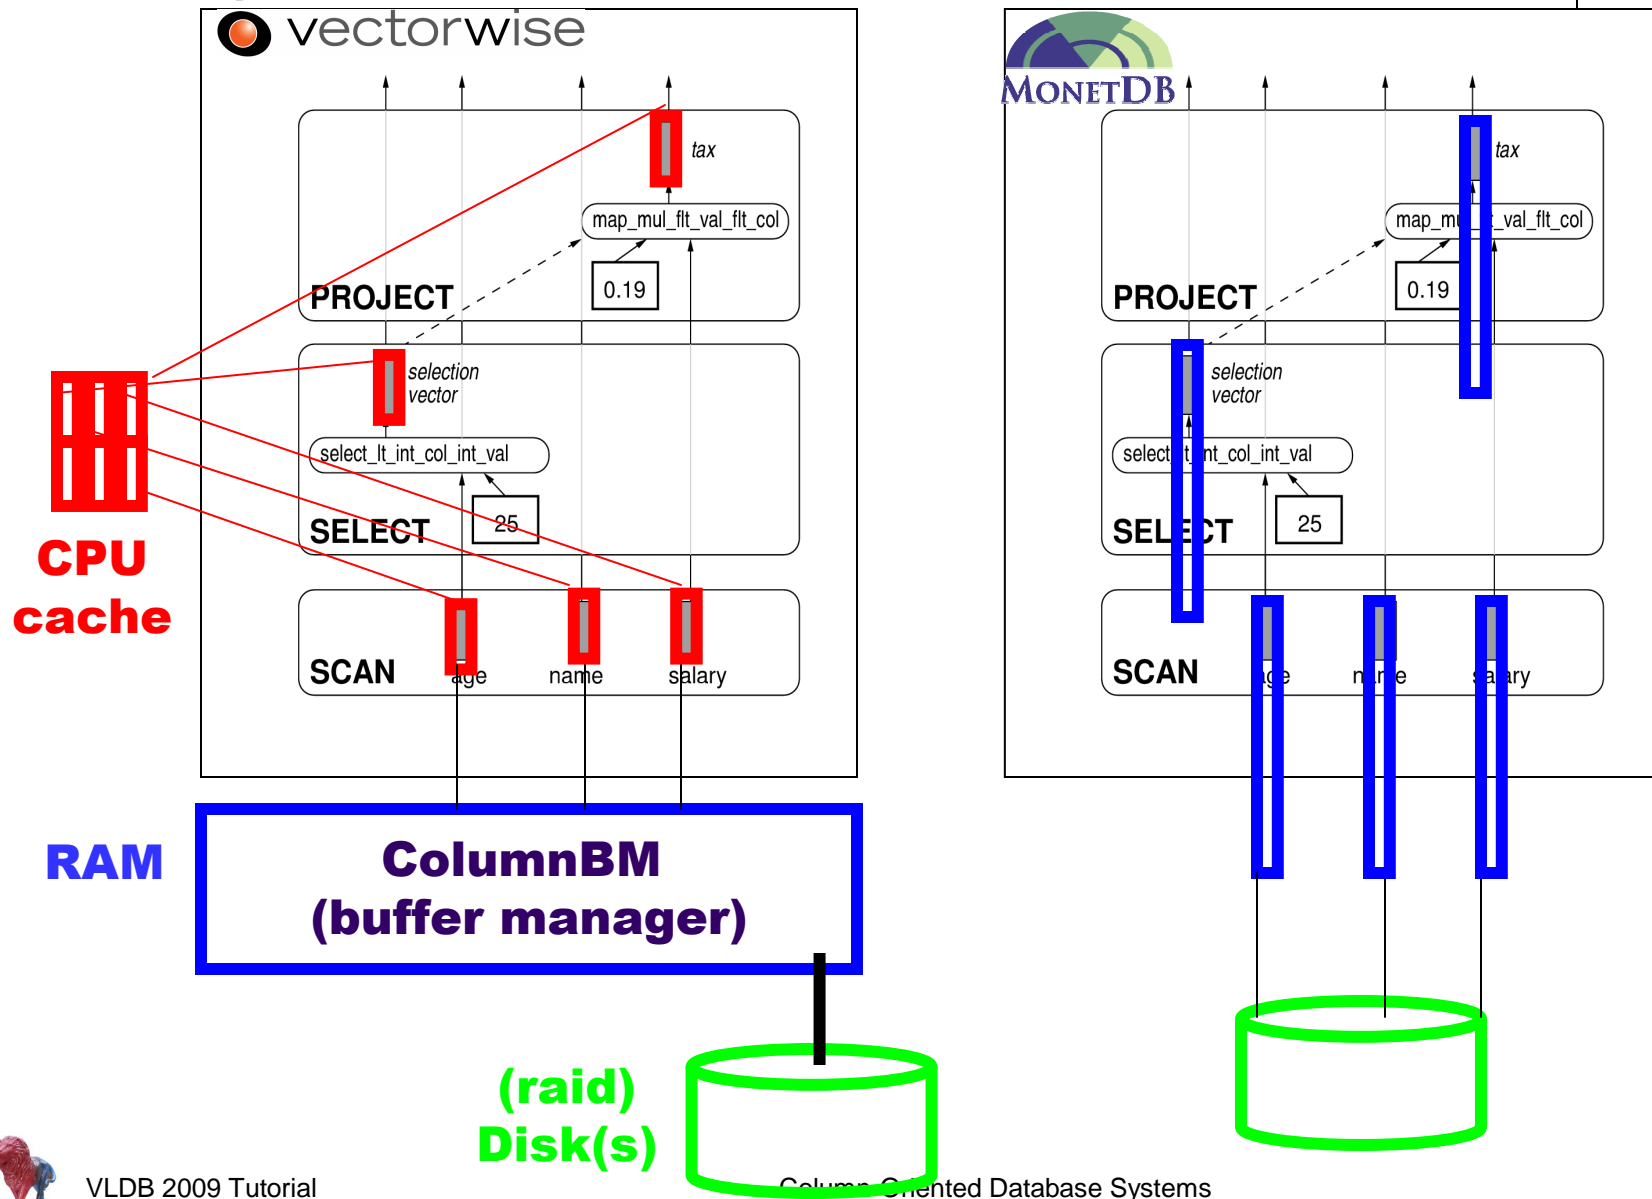


Varying the Vector size



“MonetDB/X100: Hyper-Pipelining Query Execution” Boncz, Zukowski, Nes, CIDR’05

MonetDB/MIL materializes columns





“MonetDB/X100: Hyper-Pipelining Query Execution ” Boncz, Zukowski, Nes, CIDR’05



Benefits of Vectorized Processing

1 **100x less Function Calls**

- 1 iterator.next(), primitives

1 **No Instruction Cache Misses**

- 1 High locality in the primitives

1 **Less Data Cache Misses**

- 1 Cache-conscious data placement

1 **No Tuple Navigation**

- 1 Primitives are record-oblivious, only see arrays

1 **Vectorization allows algorithmic optimization**

- 1 Move activities out of the loop (“strength reduction”)

1 **Compiler-friendly function bodies**

- 1 Loop-pipelining, automatic SIMD

“Buffering Database Operations for Enhanced Instruction Cache Performance”
Zhou, Ross, SIGMOD’04

“Block oriented processing of relational database operations in modern computer architectures”
Padmanabhan, Malkemus, Agarwal, ICDE’01





“Balancing Vectorized Query Execution with
Bandwidth Optimized Storage” Zukowski, CWI 2008



Vectorizing Relational Operators

- 1 Project
- 1 Select
 - 1 Exploit selectivities, test buffer overflow
- 1 Aggregation
 - 1 Ordered, Hashed
- 1 Sort
 - 1 Radix-sort nicely vectorizes
- 1 Join
 - 1 Merge-join + Hash-join




Column-Oriented Database Systems

VLDB
2009
Tutorial



Efficient Column Store Compression





Key Ingredients

“Super-Scalar RAM-CPU Cache Compression”
Zukowski, Heman, Nes, Boncz, ICDE’06



- 1 Compress relations on a per-column basis
 - 1 Columns compress well
- 1 Decompress small **vectors** of tuples from a column into the CPU cache
 - 1 Minimize main-memory overhead
- 1 Use light-weight, CPU-efficient algorithms
 - 1 Exploit processing power of modern CPUs





Key Ingredients

“Super-Scalar RAM-CPU Cache Compression”
Zukowski, Heman, Nes, Boncz, ICDE’06




- 1 Compress relations on a per-column basis
 - 1 Columns compress well
- 1 Decompress small **vectors** of tuples from a column into the CPU cache
 - 1 Minimize main-memory overhead





Vectorized Decompression



Idea:

decompress a vector only


compression:

- between **CPU** and **RAM**
- Instead of **disk** and **RAM** (classic)





RAM-Cache Decompression




- 1 Decompress vectors on-demand into the cache
- 1 RAM-Cache boundary only crossed once
- 1 **More (compressed) data cached in RAM**
- 1 **Less bandwidth use**





Multi-Core Bandwidth & Compression

Performance Degradation with Concurrent Queries





“Super-Scalar RAM-CPU Cache Compression”
Zukowski, Heman, Nes, Boncz, ICDE’06




CPU Efficient Decompression

- 1 Decoding loop over cache-resident vectors of code words
- 1 Avoid control dependencies within decoding loop
 - 1 no if-then-else constructs in loop body
- 1 Avoid data dependencies between loop iterations





Disk Block Layout




- 1 Forward growing section of arbitrary size **code words** (code word size fixed per block)





Disk Block Layout



- 1 Forward growing section of arbitrary size code words (code word size fixed per block)
- 1 Backwards growing **exception list**





Naïve Decompression Algorithm

- 1 Use reserved value from code word domain (MAXCODE) to *mark* exception positions


```
int code[n]; /* temporary machine addressable buffer ,  
  
/* blow up next vector of b-bit input code words into  
   machine addressable representation */  
UNPACK[b](code, input, n) ;
```

```
for(i=j=0; i<n; i++) {  
    if (code[i] < MAXCODE) {  
        output[i] = DECODE(code[i]);  
    } else {  
        output[i] = exception[--j]);  
    }  
}
```





Deterioration With Exception%




1 1.2GB/s deteriorates to 0.4GB/s





Deterioration With Exception%




1 Perf Counters: CPU mispredicts if-then-else






Patching

- 1 Maintain a *patch-list* through code word section that links exception positions





Patching



- 1 Maintain a *patch-list* through code word section that links exception positions
- 1 After decoding, *patch* up the exception positions with the correct values





Patched Decompression

```
/* initialize cur to index of first exception within codes */  
int cur = first_exception;  
int code[n]; /* temporary machine addressable buffer /
```

```
/* blow up next vector of b-bit input code words into machine  
   addressable representation */
```

```
UNPACK[b](code, input, n) ;
```

```
/* LOOP1: decode all values */  
for(int i=0; i<n; i++) {  
    output[i] = DECODE(code[i]);  
}
```

```
/* LOOP2: patch it up */  
for(int i=1; cur < n; i++) {  
    output[cur] = exception[-i];  
    cur = cur + code[cur];  
}
```





Patched Decompression

```
/* initialize cur to index of first exception within codes */
int cur = first_exception;
int code[n]; /* temporary machine addressable buffer /

/* blow up next vector of b-bit input code words into machine
   addressable representation */
UNPACK[b](code, input, n) ;

/* LOOP1: decode all values */
for(int i=0; i<n; i++) {
    output[i] = DECODE(code[i]);
}

/* LOOP2: patch it up */
for(int i=1; cur < n; i++) {
    output[cur] = exception[-i];
    cur = cur + code[cur];
}
```





Decompression Bandwidth


Xeon Decompression Bandwidth

Patching can be applied to:

- Frame of Reference (PFOR)
- Delta Compression (PFOR-DELTA)
- Dictionary Compression (PDICT)

Makes these methods much more applicable to noisy data

⌊ without additional cost



⌊ Patching makes two passes, but is faster!




Column-Oriented Database Systems

VLDB
2009
Tutorial



Conclusion





Summary (1/2)

- 1 Columns and Row-Stores: different?
 - 1 No fundamental differences
 - 1 Can current row-stores simulate column-stores now?
 - 1 not efficiently: row-stores need change
 - 1 On disk layout vs execution layout
 - 1 actually independent issues, on-the-fly conversion pays off
 - 1 column favors sequential access, row random
 - 1 Mixed Layout schemes
 - 1 Fractured mirrors
 - 1 PAX, Clotho
 - 1 Data morphing





Summary (2/2)

1 Crucial Columnar Techniques

1 Storage

- 1 Lean headers, sparse indices, fast positional access

1 Compression

- 1 Operating on compressed data
- 1 Lightweight, vectorized decompression

1 Late vs Early materialization

- 1 Non-join: LM always wins
- 1 Naïve/Invisible/Jive/Flash/Radix Join (LM often wins)

1 Execution

- 1 Vectorized in-cache execution
- 1 Exploiting SIMD





Future Work

- 1 looking at write/load tradeoffs in column-stores
 - 1 read-only vs batch loads vs trickle updates vs OLTP





Updates (1/3)

- 1 Column-stores are update-in-place averse
 - 1 In-place: I/O for each column
 - 1 + re-compression
 - 1 + multiple sorted replicas
 - 1 + sparse tree indices

Update-in-place is infeasible!





Updates (2/3)

- 1 Column-stores use differential mechanisms instead
 - 1 Differential lists/files or more advanced (e.g. PDTs)
 - 1 Updates buffered in RAM, merged on each query
 - 1 Checkpointing merges differences in bulk sequentially
 - 1 I/O trends favor this anyway
 - § trade RAM for converting random into sequential I/O
 - § this trade is also needed in Flash (do not write randomly!)
 - 1 How high loads can it sustain?
 - § Depends on available RAM for buffering (how long until full)
 - § Checkpoint must be done within that time
 - § The longer it can run, the less it molests queries
 - § Using Flash for buffering differences buys a lot of time
 - § Hundreds of GBs of differences per server





Updates (3/3)

- 1 Differential transactions favored by hardware trends
- 1 Snapshot semantics accepted by the user community
 - 1 can always convert to serialized
 - “Serializable Isolation For Snapshot Databases” Alomari, Cahill, Fekete, Roehm, SIGMOD’08
- ⌚ Row stores could also use differential transactions and be efficient!
 - ⌚ Implies a departure from ARIES
 - ⌚ Implies a full rewrite

My conclusion:

a system that combines row- and columns needs differentially implemented transactions.

Starting from a pure column-store, this is a limited add-on.

Starting from a pure row-store, this implies a full rewrite.





Future Work

- 1 looking at write/load tradeoffs in column-stores
 - 1 read-only vs batch loads vs trickle updates vs OLTP
- 1 database design for column-stores
- 1 column-store specific optimizers
 - 1 compression/materialization/join tricks 1 cost models?
- 1 hybrid column-row systems
 - 1 can row-stores learn new column tricks?
 - 1 Study of the minimal number changes one needs to make to a row store to get the majority of the benefits of a column-store
 - 1 Alternative: add features to column-stores that make them more like row stores





Conclusion

- 1 Columnar techniques provide clear benefits for:
 - 1 Data warehousing, BI
 - 1 Information retrieval, graphs, e-science
- 1 A number of crucial techniques make them effective
 - 1 Without these, existing row systems do not benefit
- 1 Row-Stores and column-stores could be combined
 - 1 Row-stores may adopt some column-store techniques
 - 1 Column-stores add row-store (or PAX) functionality
- 1 Many open issues to do research on!

